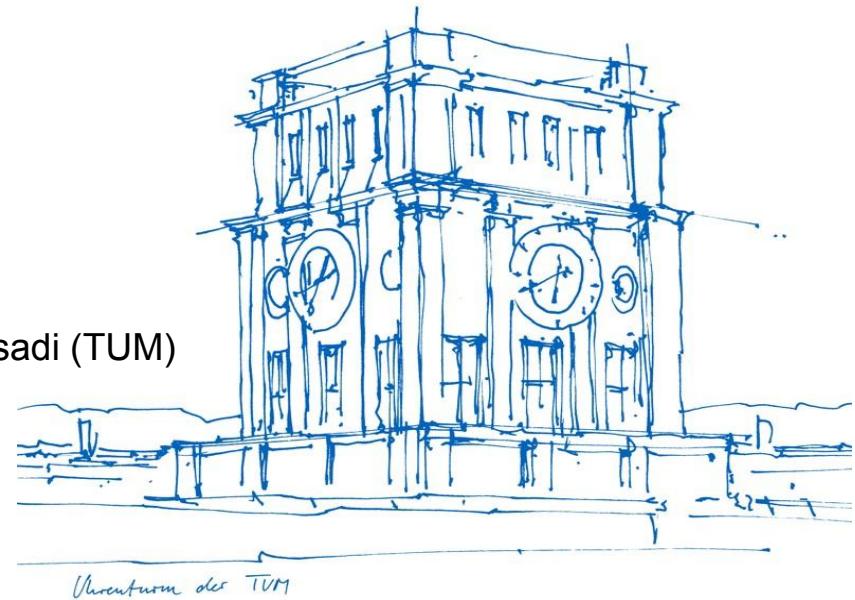


VM Selection for Financial Exchanges in the Cloud

Master Thesis Final Presentation

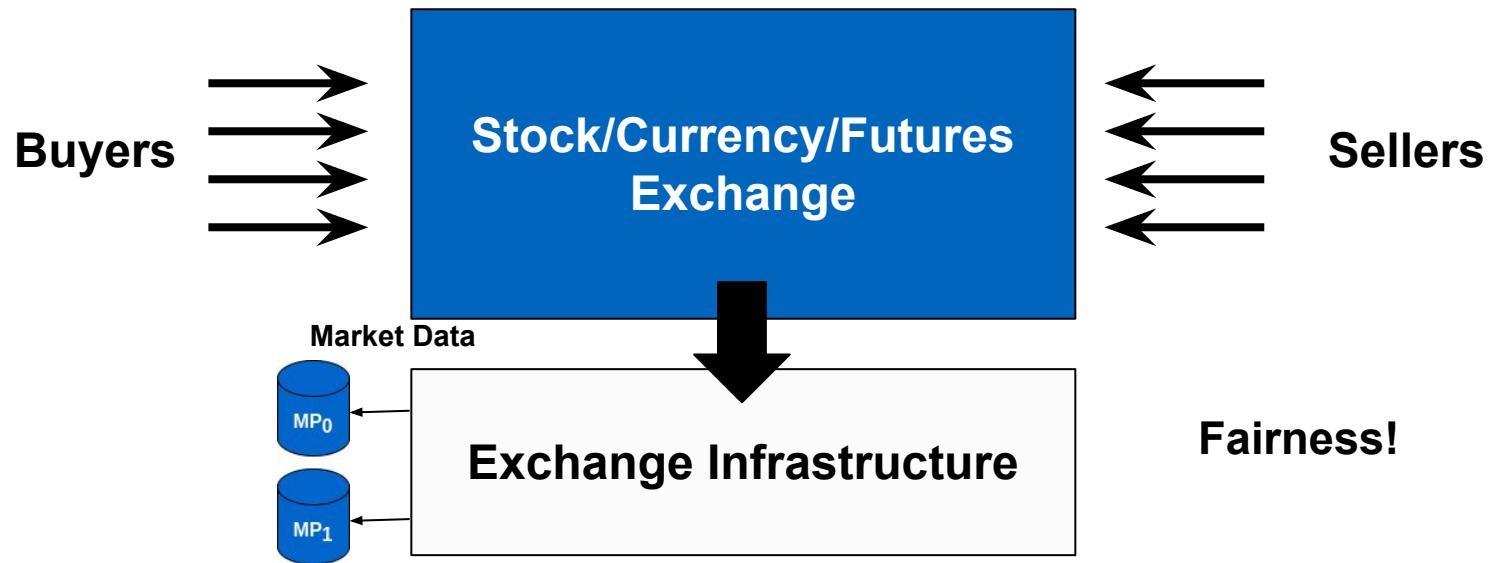
Student: Daniel Duclos-Cavalcanti

Supervisors: Muhammad Haseeb (NYU), Navidreza Asadi (TUM)



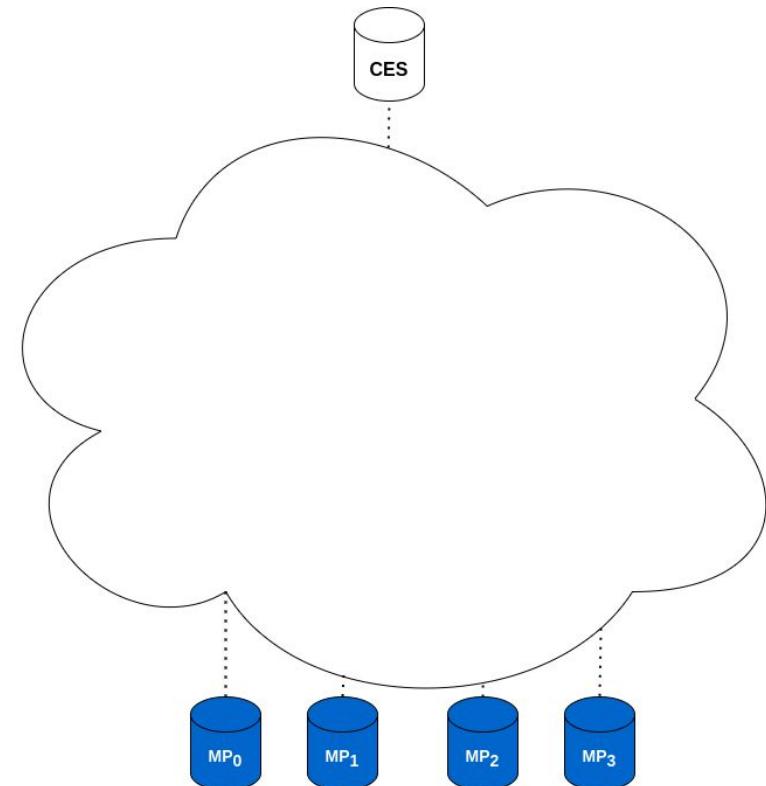
Background: Financial Exchanges and the Cloud

- What is a Financial Exchange?
- Migration: on-premise => public cloud
- **Cloud Medium:** i) High Variance ii) ~~Mechanisms~~ for fair delivery



State-of-the-Art: Jasper¹: Scalable Exchange in the Cloud

- Scalable, Fair, and Performant Multicast
- Market Data Dissemination

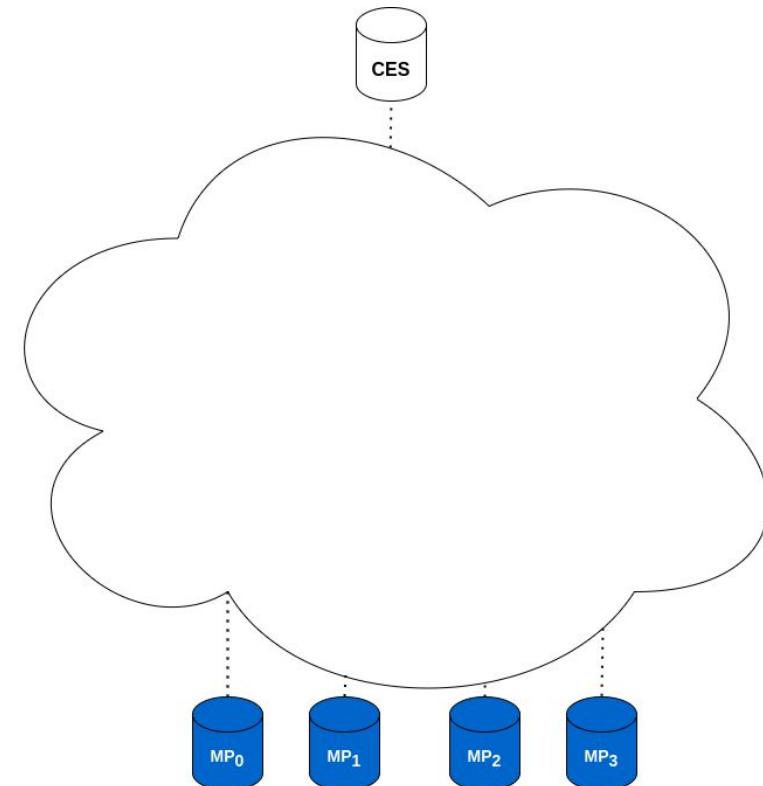


¹Jasper: <https://arxiv.org/abs/2402.09527>, <https://sigcomm24posterdemo.hotcrp.com/paper/>

²Hyugens: <https://www.usenix.org/system/files/conference/nsdi18/nsdi18-geng.pdf>

State-of-the-Art: Jasper¹: Scalable Exchange in the Cloud

- Scalable, Fair, and Performant Multicast
- Market Data Dissemination
- Techniques:

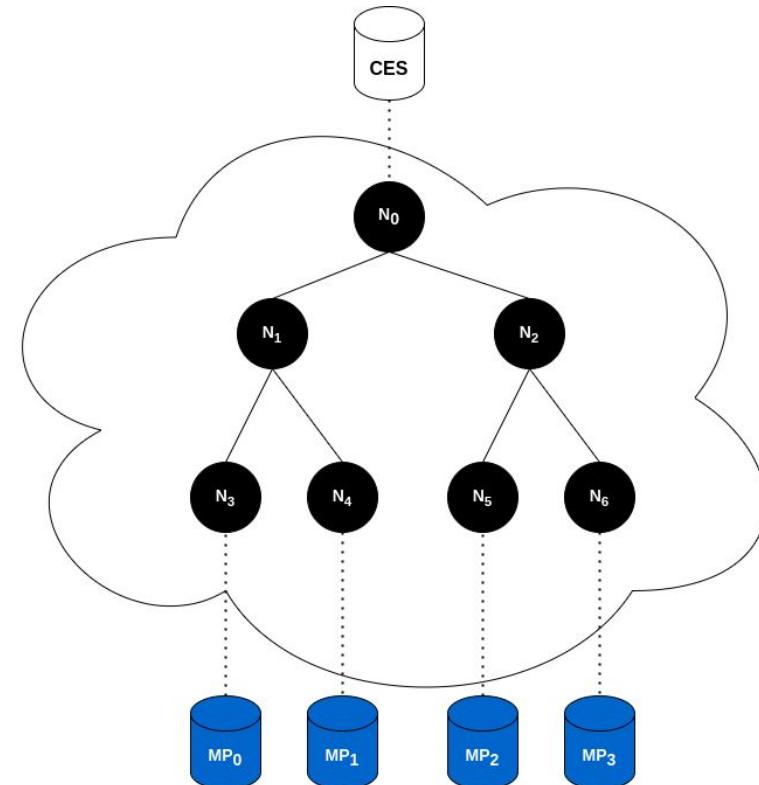


¹Jasper: <https://arxiv.org/abs/2402.09527>, <https://sigcomm24posterdemo.hotcrp.com/paper/>

²Hyugens: <https://www.usenix.org/system/files/conference/nsdi18/nsdi18-geng.pdf>

State-of-the-Art: Jasper¹: Scalable Exchange in the Cloud

- Scalable, Fair, and Performant Multicast
- Market Data Dissemination
- Techniques:
 - Overlay Multicast Tree

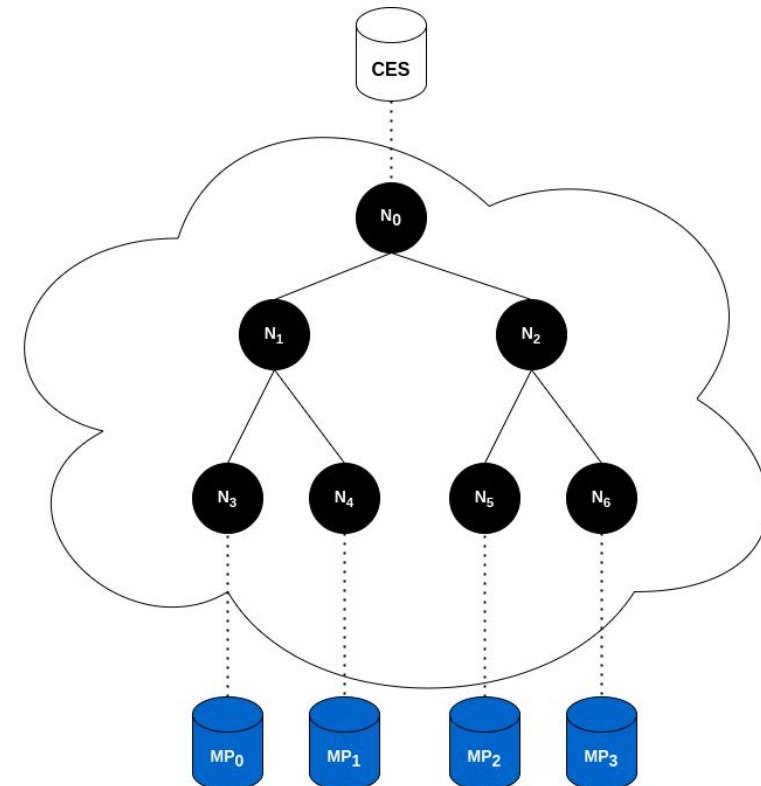


¹Jasper: <https://arxiv.org/abs/2402.09527>, <https://sigcomm24posterdemo.hotcrp.com/paper/>

²Hyugens: <https://www.usenix.org/system/files/conference/nsdi18/nsdi18-geng.pdf>

State-of-the-Art: Jasper¹: Scalable Exchange in the Cloud

- Scalable, Fair, and Performant Multicast
- Market Data Dissemination
- Techniques:
 - Overlay Multicast Tree
 - Clock-Synchronization²

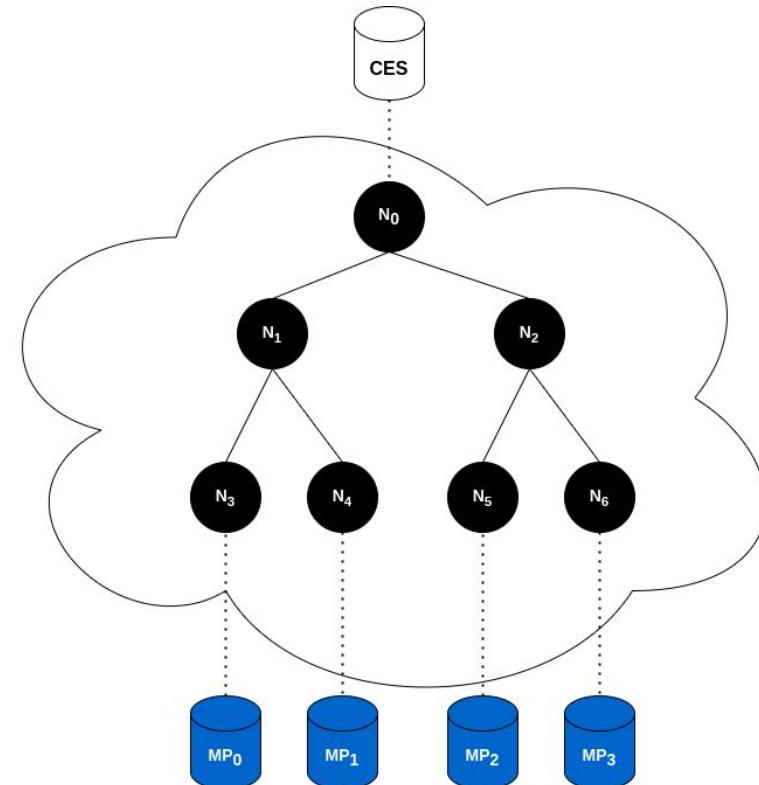


¹Jasper: <https://arxiv.org/abs/2402.09527>, <https://sigcomm24posterdemo.hotcrp.com/paper/>

²Hyugens: <https://www.usenix.org/system/files/conference/nsdi18/nsdi18-geng.pdf>

State-of-the-Art: Jasper¹: Scalable Exchange in the Cloud

- Scalable, Fair, and Performant Multicast
- Market Data Dissemination
- Techniques:
 - Overlay Multicast Tree
 - Clock-Synchronization²
 - Hold-and-Release



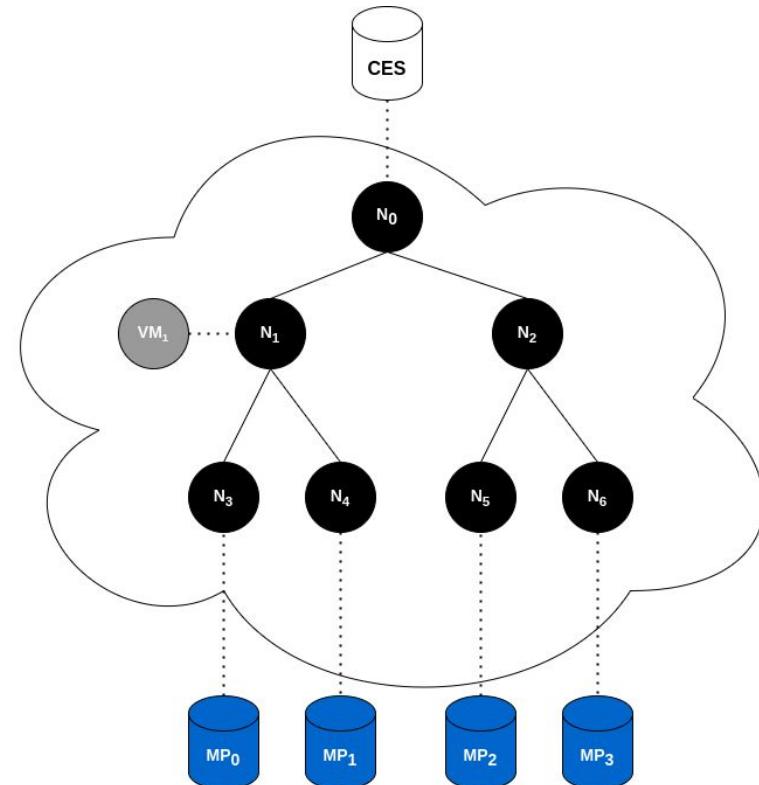
¹Jasper: <https://arxiv.org/abs/2402.09527>, <https://sigcomm24posterdemo.hotcrp.com/paper/>

²Hyugens: <https://www.usenix.org/system/files/conference/nsdi18/nsdi18-geng.pdf>

State-of-the-Art: Jasper¹: Scalable Exchange in the Cloud

- Scalable, Fair, and Performant Multicast
- Market Data Dissemination
- Techniques:
 - Overlay Multicast Tree
 - Clock-Synchronization²
 - Hold-and-Release
- Opportunity:
 - “Stragglers”
 - Underperforming VMs

Stragglers affect deadlines!



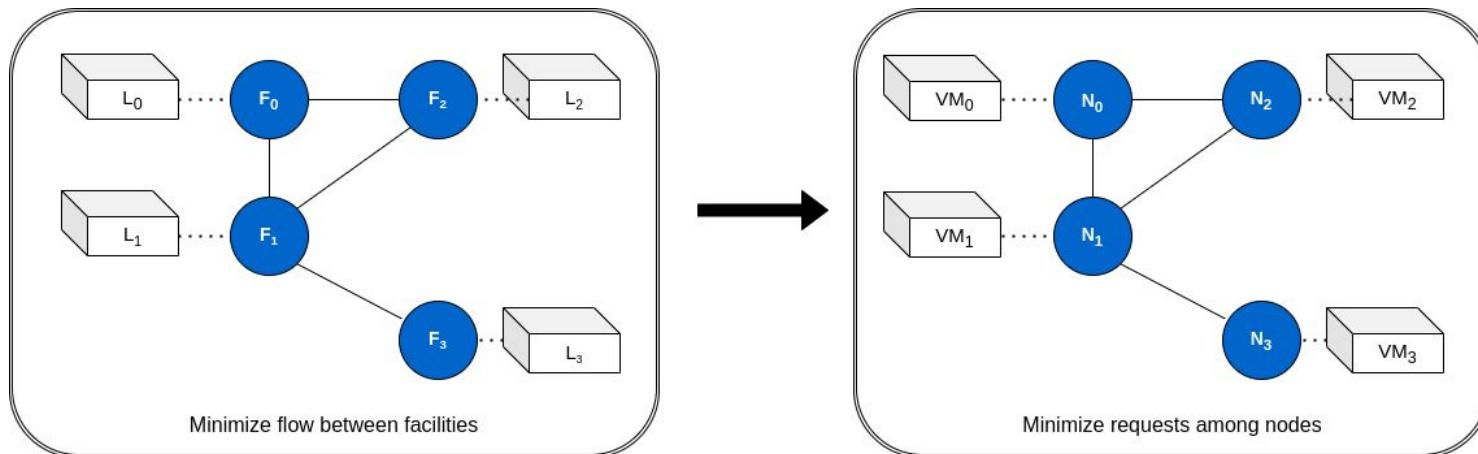
¹Jasper: <https://arxiv.org/abs/2402.09527>, <https://sigcomm24posterdemo.hotcrp.com/paper/>

²Hyugens: <https://www.usenix.org/system/files/conference/nsdi18/nsdi18-geng.pdf>

State-of-the-Art: LemonDrop¹

- Method of Selecting/Scheduling VMs on a cluster
- Framed QAP => NP-Hard => Scalability issues

Disproportionally Robust for Jasper's use-case



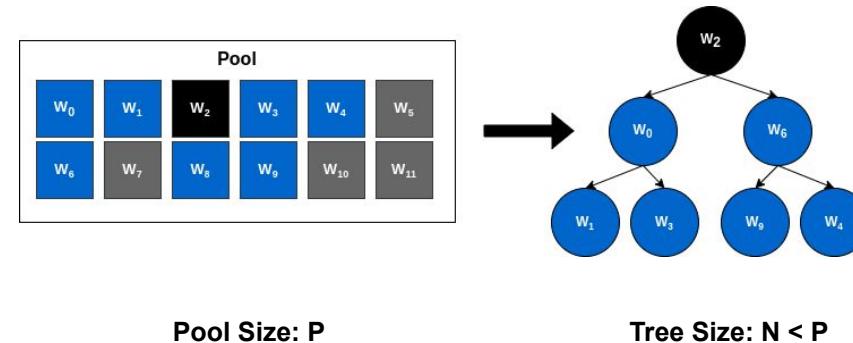
¹LemonDrop: <https://searchworks.stanford.edu/view/14423035>

Develop lightweight VM selection heuristic

- Minimize “Straggler” Effect
- Lightweight and Scalable

TreeBuilder¹: Testbench and Design

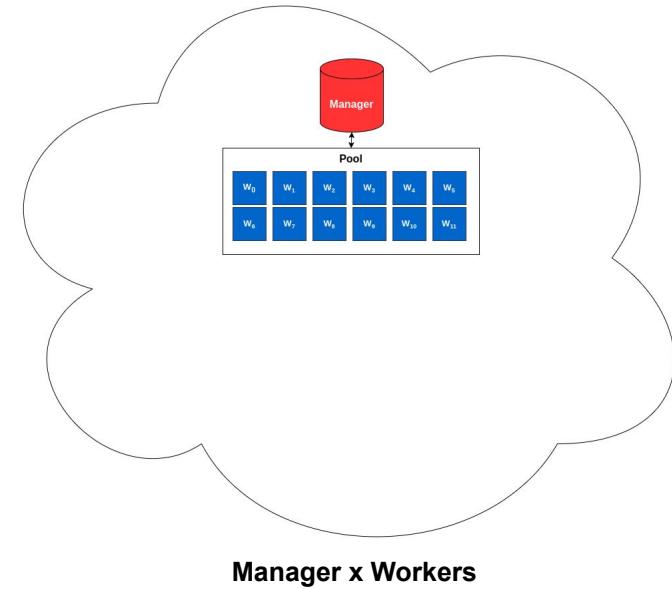
- **Goal:** VM Selection Method to minimize “stragglers”
- **Testbench:**



¹TreeBuilder: <https://github.com/duclos-cavalcanti/TreeBuilder>

TreeBuilder¹: Testbench and Design

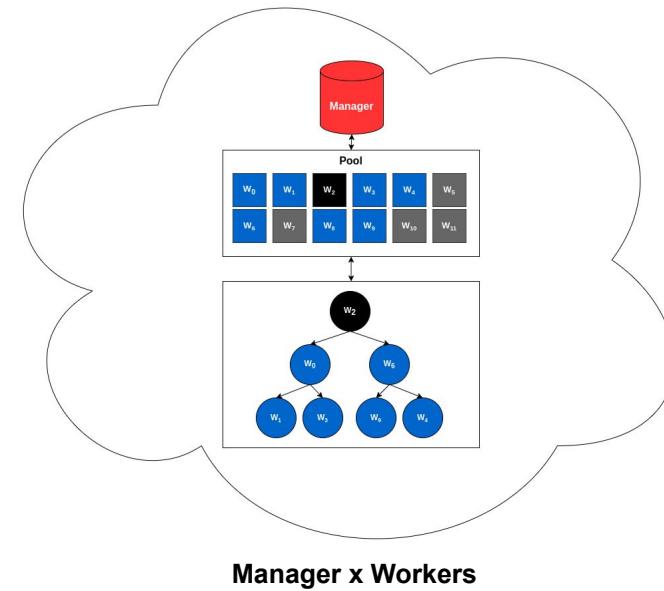
- **Goal:** VM Selection Method to minimize “stragglers”
- **Testbench:**
 1. Launch a cluster of P VMs
 - **Workers => Pool**
 - **Manager => Supervisor**



¹TreeBuilder: <https://github.com/duclos-cavalcanti/TreeBuilder>

TreeBuilder¹: Testbench and Design

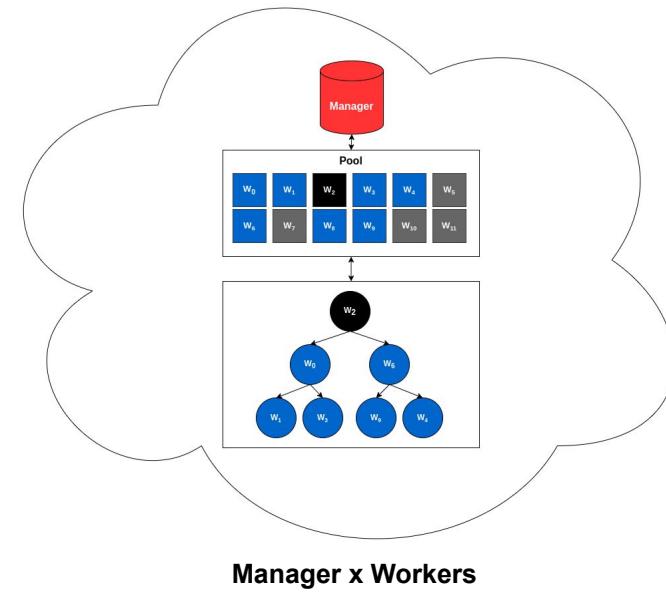
- **Goal:** VM Selection Method to minimize “stragglers”
- **Testbench:**
 1. Launch a cluster of P VMs
 2. Apply Techniques:



¹TreeBuilder: <https://github.com/duclos-cavalcanti/TreeBuilder>

TreeBuilder¹: Testbench and Design

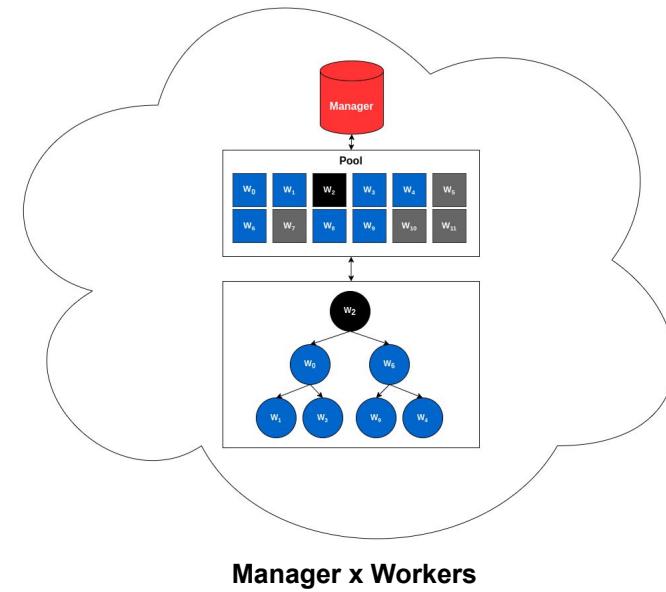
- **Goal:** VM Selection Method to minimize “stragglers”
- **Testbench:**
 1. Launch a cluster of P VMs
 2. Apply Techniques:
 - Proposed Method



¹TreeBuilder: <https://github.com/duclos-cavalcanti/TreeBuilder>

TreeBuilder¹: Testbench and Design

- **Goal:** VM Selection Method to minimize “stragglers”
- **Testbench:**
 1. Launch a cluster of P VMs
 2. Apply Techniques:
 - Proposed Method
 - LemonDrop



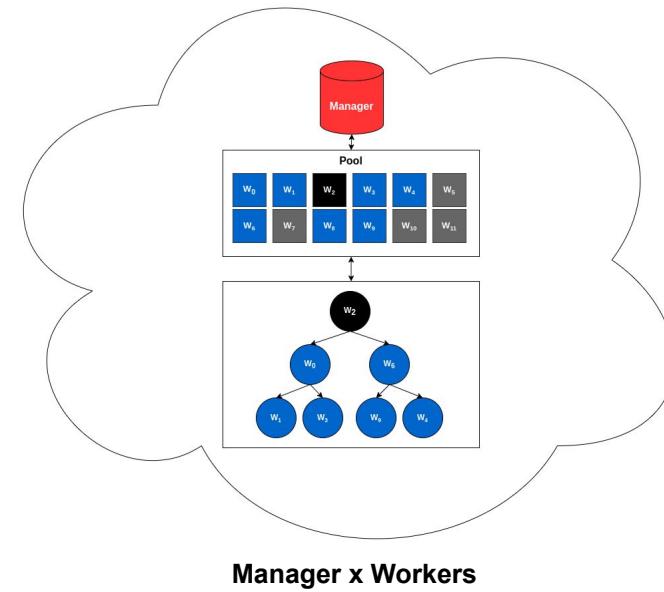
¹TreeBuilder: <https://github.com/duclos-cavalcanti/TreeBuilder>

TreeBuilder¹: Testbench and Design

- **Goal:** VM Selection Method to minimize “stragglers”

- **Testbench:**

1. Launch a cluster of P VMs
2. Apply Techniques:
 - Proposed Method
 - LemonDrop
 - Vanilla



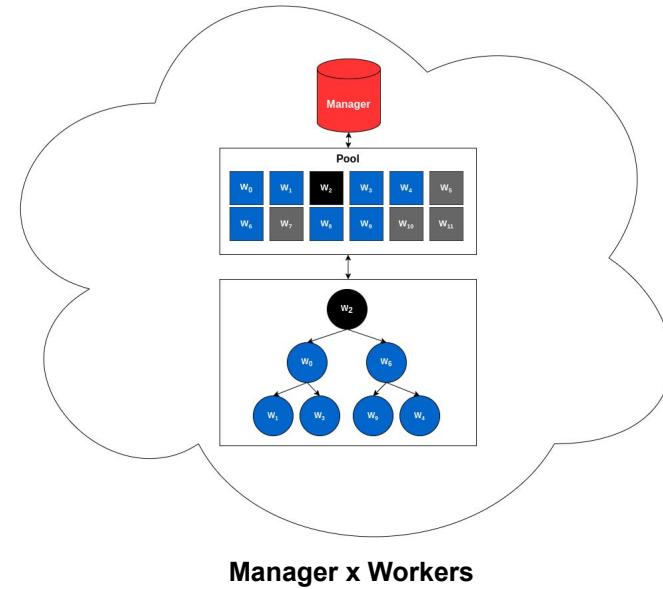
¹TreeBuilder: <https://github.com/duclos-cavalcanti/TreeBuilder>

TreeBuilder¹: Testbench and Design

- **Goal:** VM Selection Method to minimize “stragglers”
- **Testbench:**
 1. Launch a cluster of P VMs
 2. Apply Techniques:
 - Proposed Method
 - LemonDrop
 - Vanilla
 3. Evaluate trees

Q1: How are nodes examined and selected?

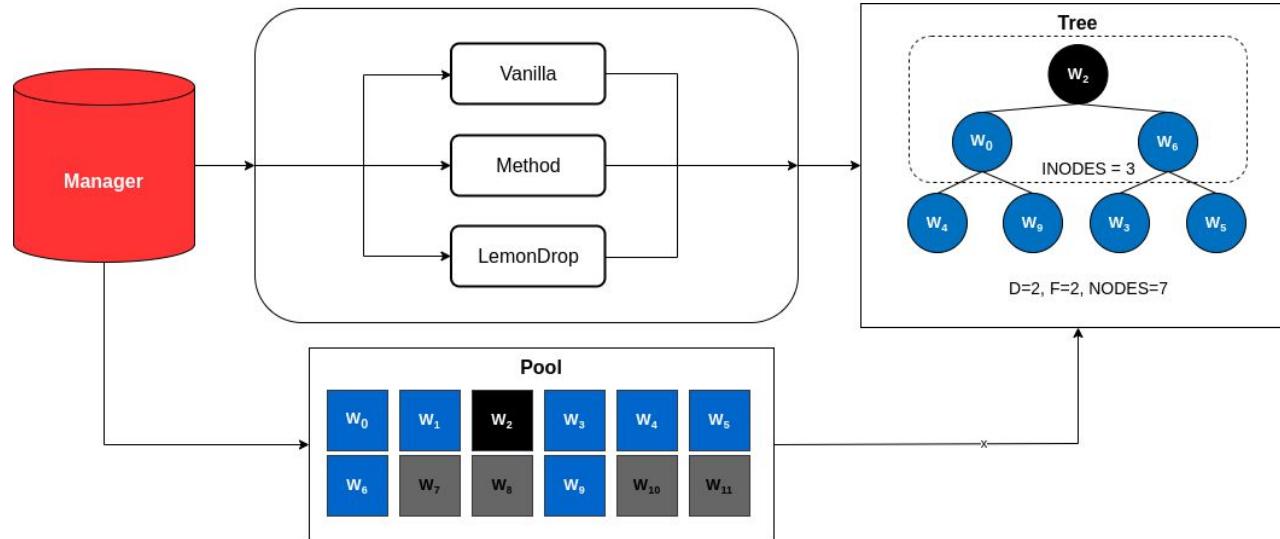
Q2: How are trees evaluated?



¹TreeBuilder: <https://github.com/duclos-cavalcanti/TreeBuilder>

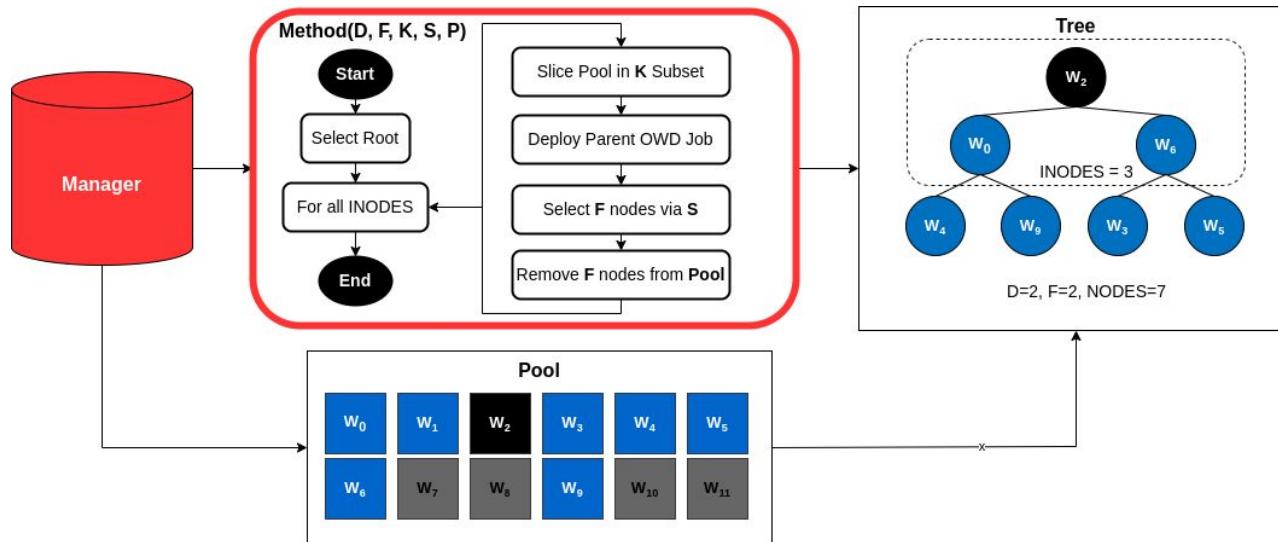
TreeBuilder: Testbench and Design

1. Launches cluster of **N** VMs (Manager x Workers)
2. Apply: [Proposed Method, LemonDrop and Vanilla]
3. Evaluate Trees



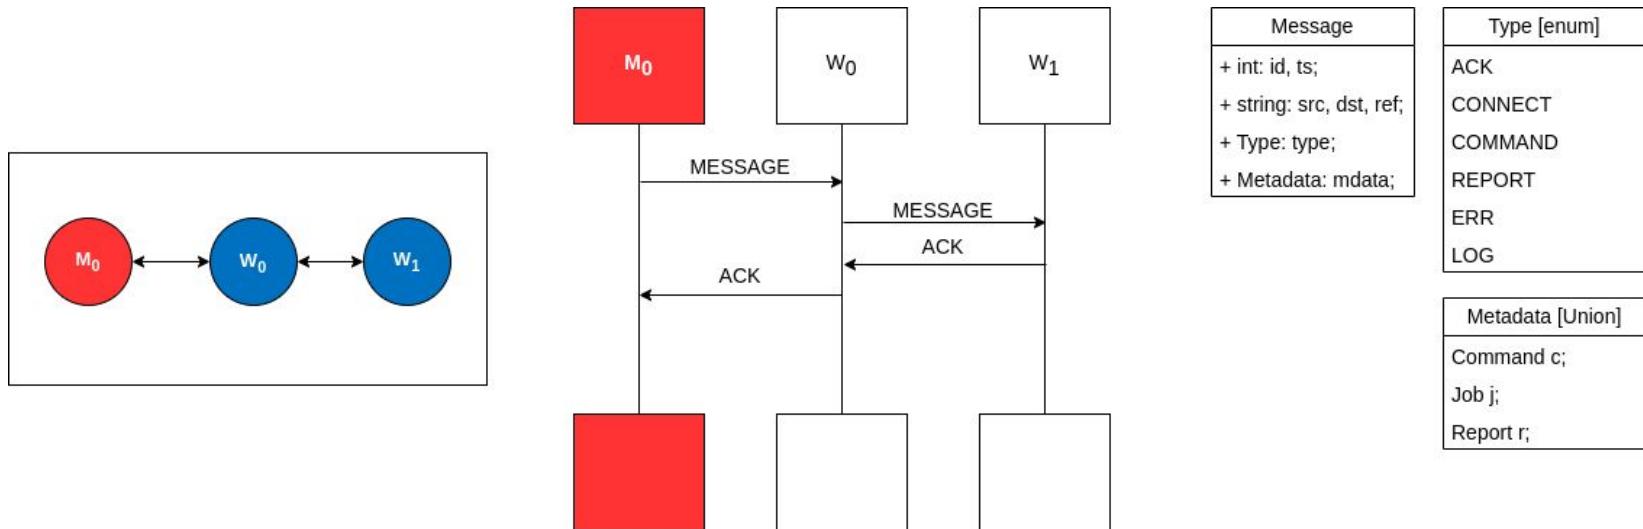
TreeBuilder: Testbench and Design

1. Launches cluster of N VMs (Manager x Workers)
2. Apply: [**Proposed Method**, LemonDrop and Vanilla]
3. Evaluate Trees



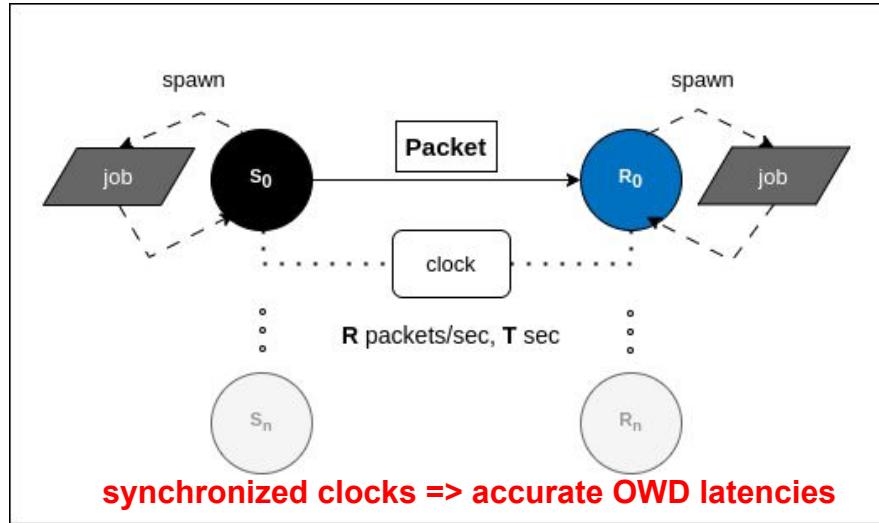
Implementation: Messages

- Request and Reply
- Defined and Flexible Data Structures



Implementation: Node(s) Examinations

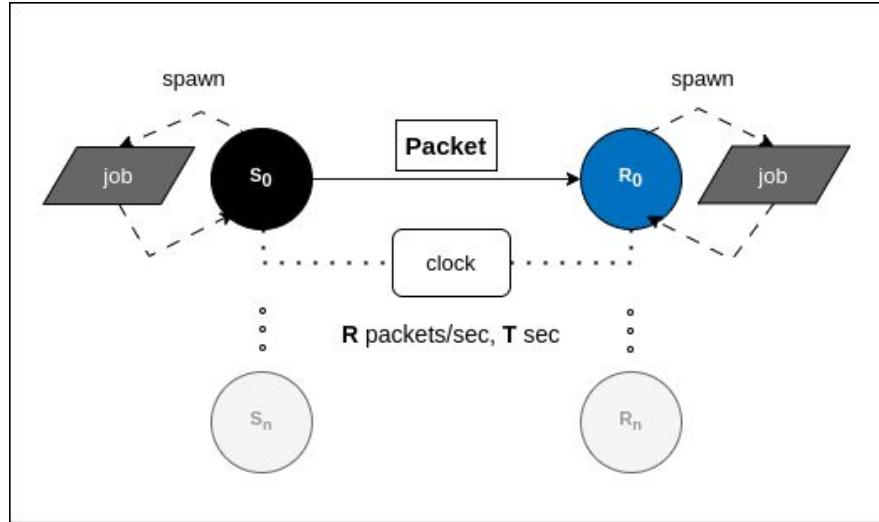
- Collective Tasks (Jobs)



The diagram shows a stack of three rectangular boxes. The top box is black and labeled "Results [R₀]". The middle box is white and contains the text "int rate, duration, recv; float p90, p75, p50, p25; float mean, stddev;". The bottom box is white and contains the text "float p90, p75, p50, p25; float mean, stddev;".

Implementation: Node(s) Examinations

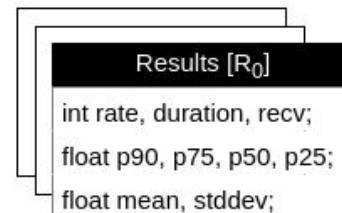
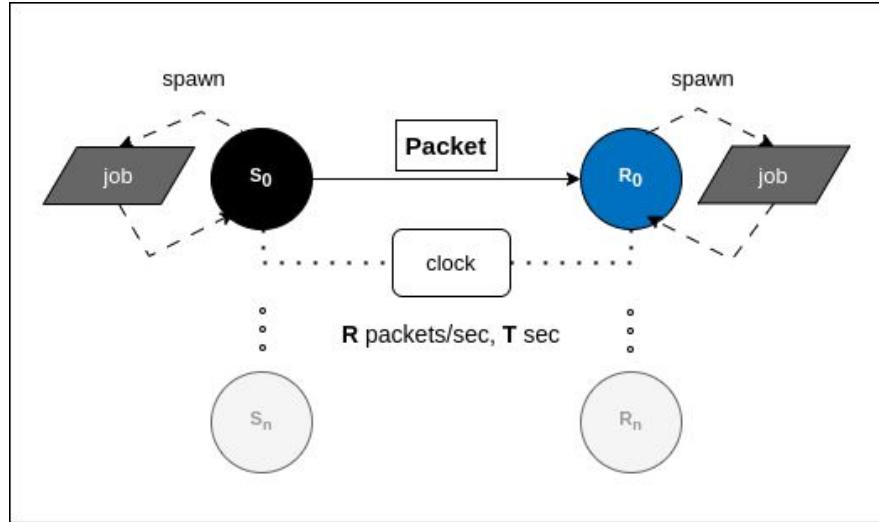
- Collective Tasks (Jobs)
- Generate Results



Results [R_0]
int rate, duration, recv;
float p90, p75, p50, p25;
float mean, stddev;

Implementation: Node(s) Examinations

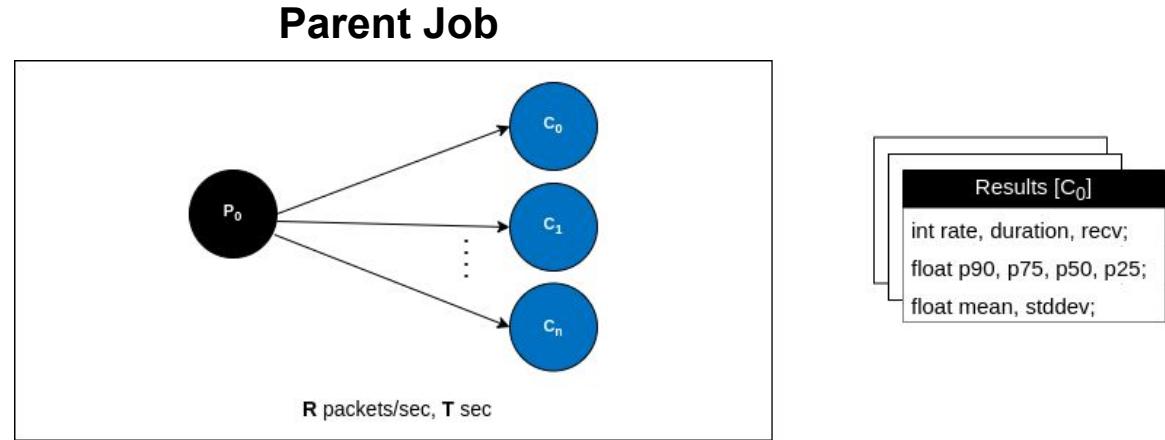
- Collective Tasks (Jobs)
- Generate Results
- Types:



Results [R_0]
int rate, duration, recv;
float p90, p75, p50, p25;
float mean, stddev;

Implementation: Node(s) Examinations

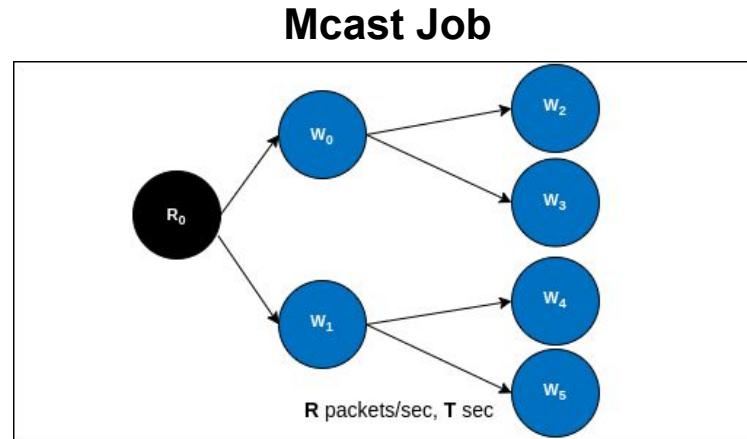
- Collective Tasks (Jobs)
- Generate Results
- Types:
 - Parent



Examines potential children nodes to a given parent!

Implementation: Node(s) Examinations

- Collective Tasks (Jobs)
- Generate Results
- Types:
 - Parent
 - Mcast

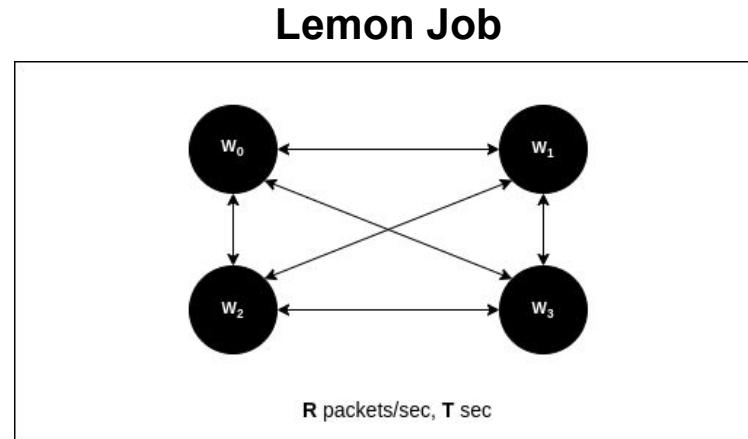


```
Results [W2]
int rate, duration, recv;
float p90, p75, p50, p25;
float mean, stddev;
```

Evaluates performance of a built tree!

Implementation: Node(s) Examinations

- Collective Tasks (Jobs)
- Generate Results
- Types:
 - Parent
 - Mcast
 - Lemon

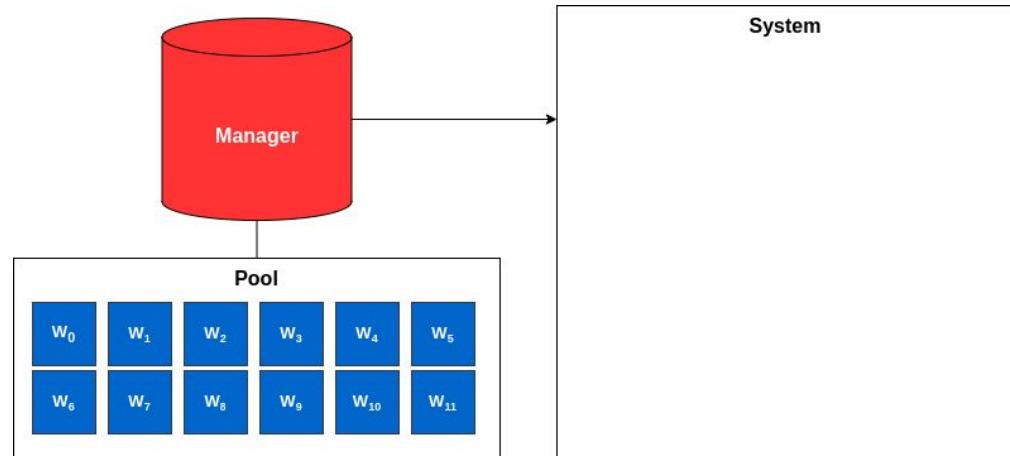
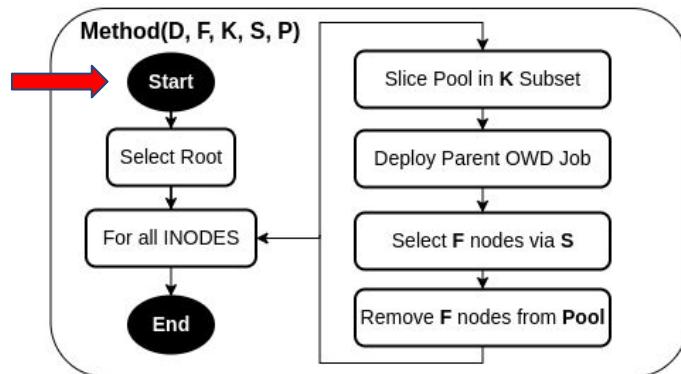


```
Results [W0]
int rate, duration, recv;
float p90, p75, p50, p25;
float mean, stddev;
```

Collects pair-wise median OWDs for LemonDrop's Heuristic!

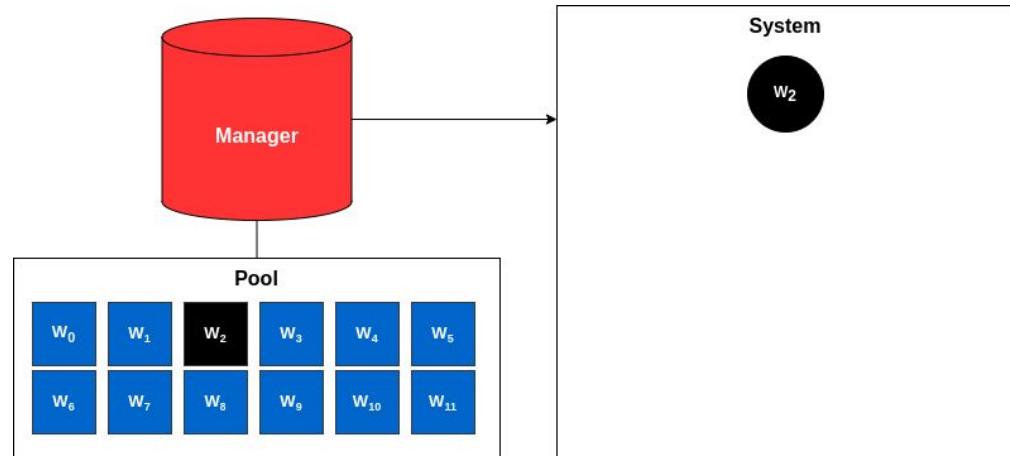
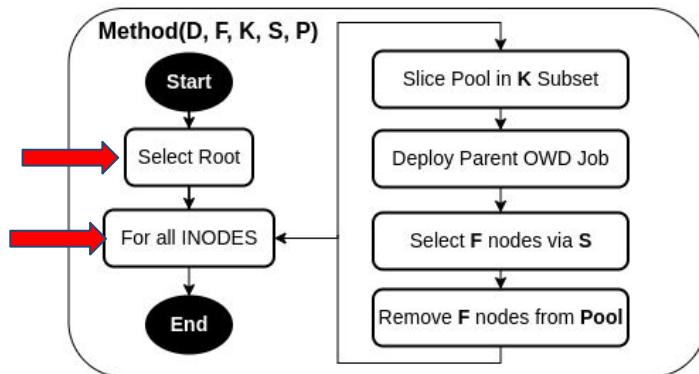
Implementation: Proposed Selection Method

- Define depth and fan-out: D , F
- Bounding parameter: K
- Strategy: S
- Pool: P



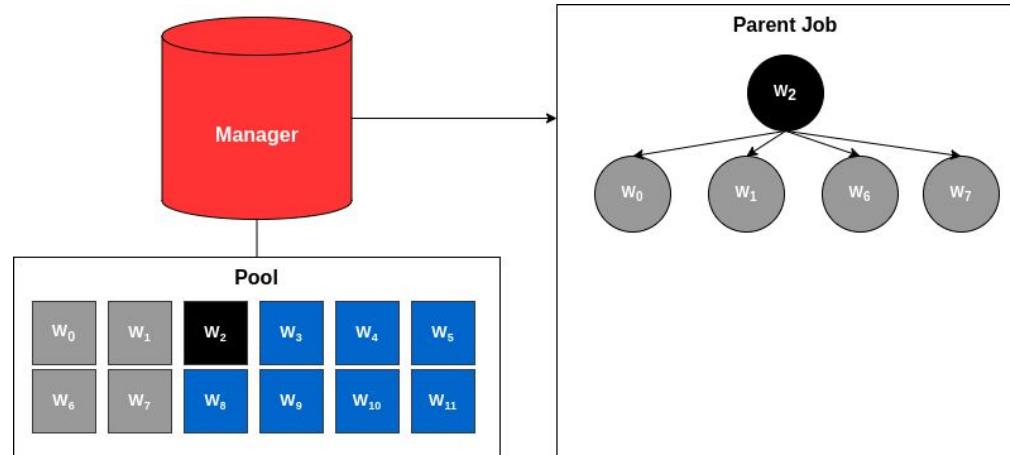
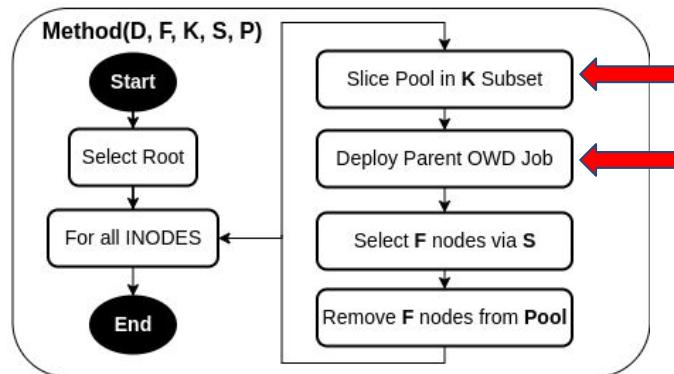
Implementation: Proposed Selection Method

- Define depth and fan-out: D , F
- Bounding parameter: K
- Strategy: S
- Pool: P



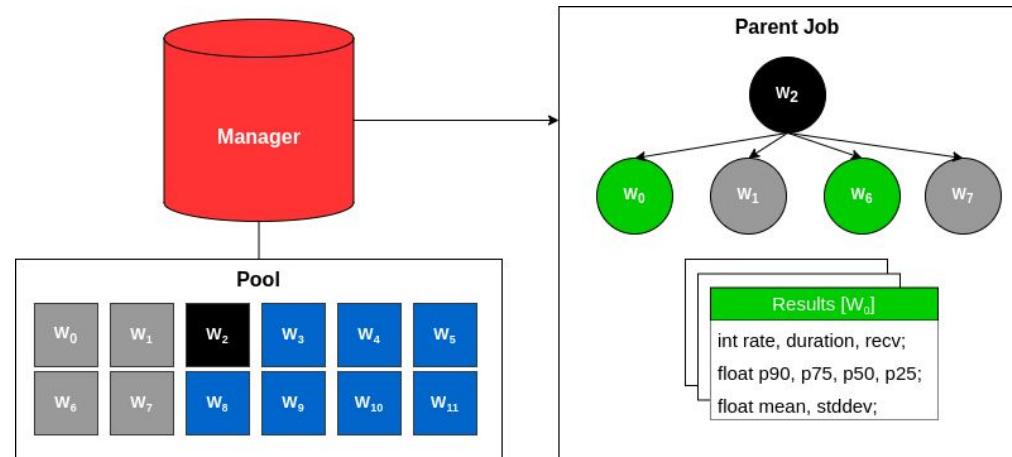
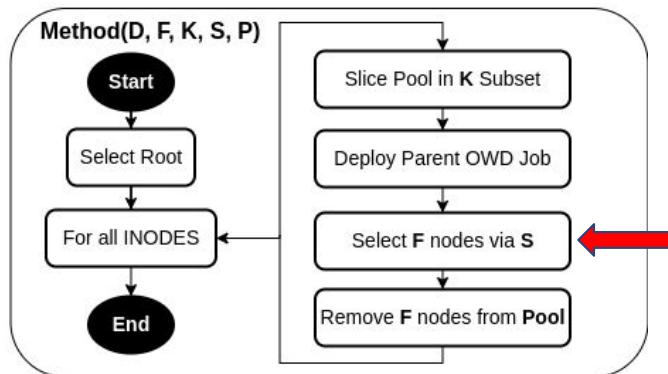
Implementation: Proposed Selection Method

- Define depth and fan-out: D , F
- Bounding parameter: K
- Strategy: S
- Pool: P



Implementation: Proposed Selection Method

- Define depth and fan-out: **D**, **F**
- Bounding parameter: **K**
- Strategy: **S**
- Pool: **P**



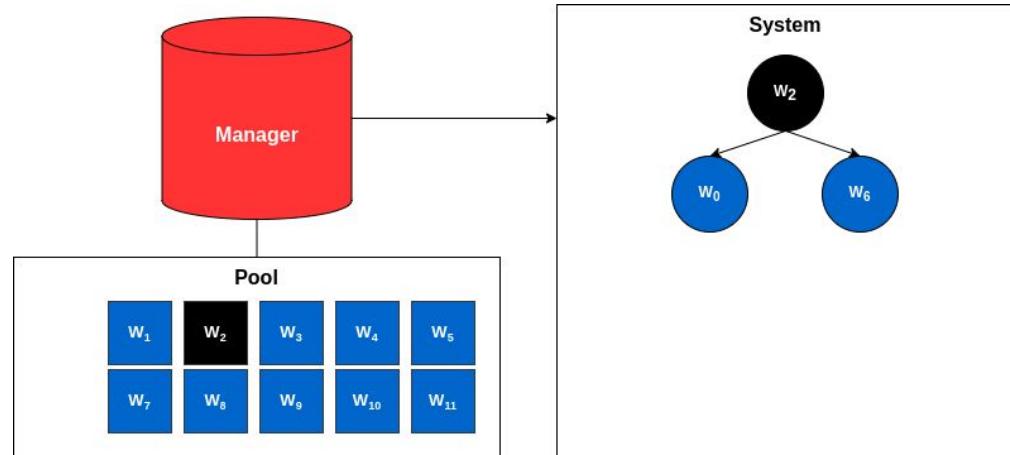
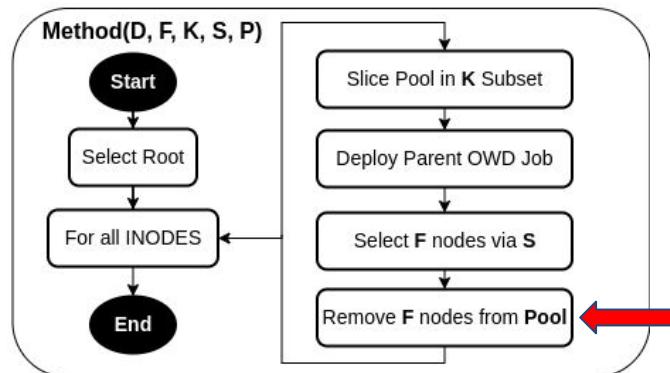
Strategies [S]:

- **P90:** $(1 \times p90)$
- **P50:** $(1 \times p50)$
- **Weighted:** $(0.7 \times \text{stddev}) + (0.3 \times p90)$

↓ lower => better ↑

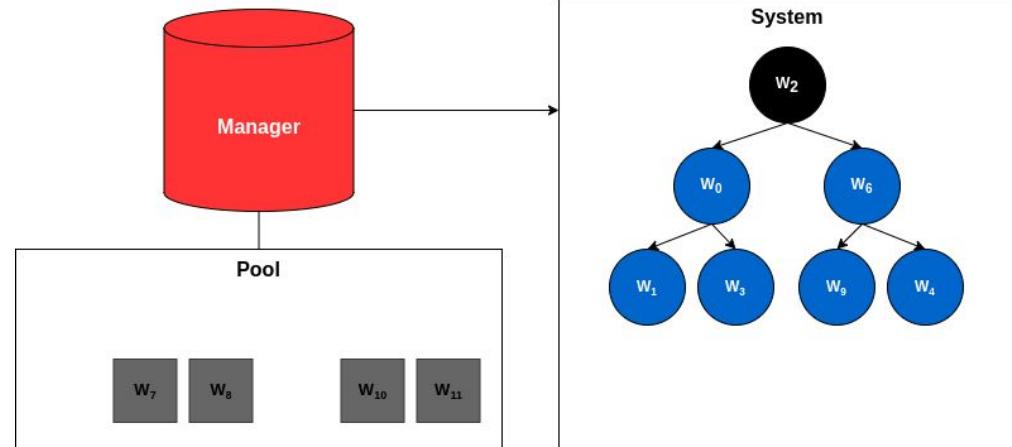
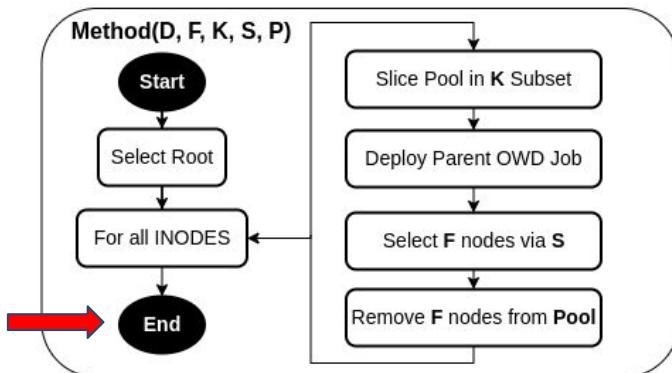
Implementation: Proposed Selection Method

- Define depth and fan-out: D , F
- Bounding parameter: K
- Strategy: S
- Pool: P



Implementation: Proposed Selection Method

- Define depth and fan-out: D , F
- Bounding parameter: K
- Strategy: S
- Pool: P



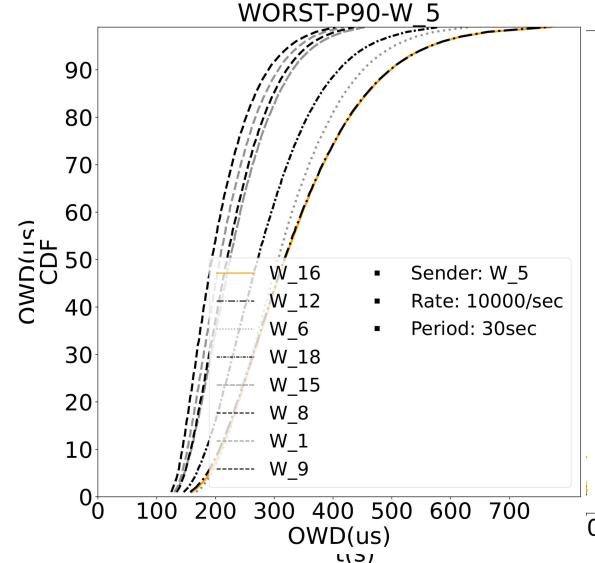
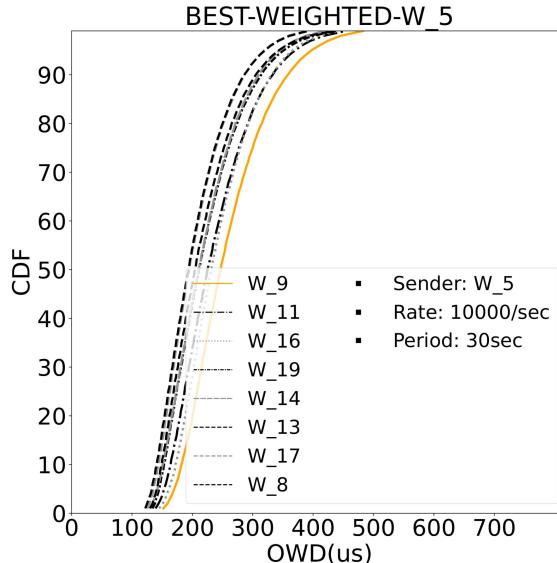
Experimental Setup

- Build Trees:
 - **Best:** P90 x P50 x WEIGHTED
 - **Worst:** P90
 - **Random**
 - **Lemon:** $\epsilon = 1 \times 10^{-4}$, $n = 1000$
- GCP Cloud Instances: “c2d-highcpu-8”
 - 8 virtual AMD Milan CPU's
 - 16GB of Memory
 - 16Gbps of Network Bandwidth

D (Depth)	3
F (Fanout)	2
K (Pool Subset)	$2^*F = 4$
P (Pool Size)	25
Duration Parent Job	10 sec
Duration Mcast Job	30 sec
Rate	10K
Number of Runs	2
Number of Roots	2

Results: Straggler Impact – BEST x WORST

- **WORST** tree displays sparser CDF profiles and right-shifted vs **BEST-WEIGHTED**
- **WORST** shows higher variancy
- **BEST-WEIGHTED** 28.11% lower OWD latency at 90th %

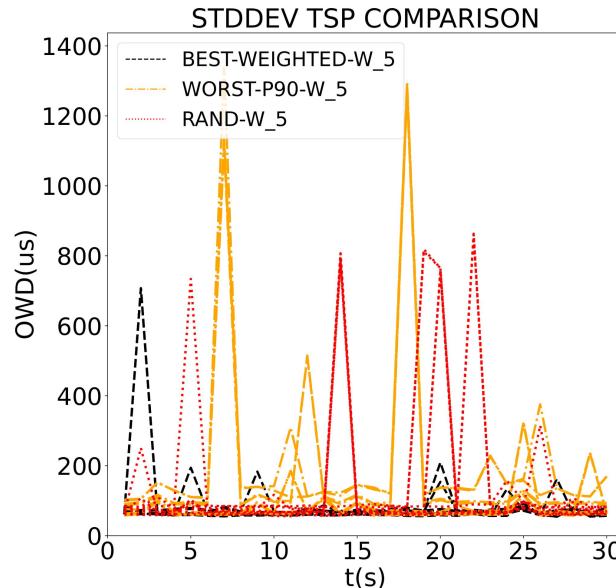


	BEST-WEIGHTED	WORST
90th %	358 μ s	498 μ s
median	247 μ s	318 μ s
mean	261.4 μ s	353.66 μ s
stddev	78.95 μ s	352.29 μ s
ID	W_9	W_16

➤ Best-WEIGHTED is uniformly better & more stable!

Results: Straggler Impact – BEST x WORST x RAND

- **BEST** 6.53% lower OWD 90th % vs **RAND**
- **RAND** in between **BEST** and **WORST**
- **RAND** displays considerable variance vs **BEST**

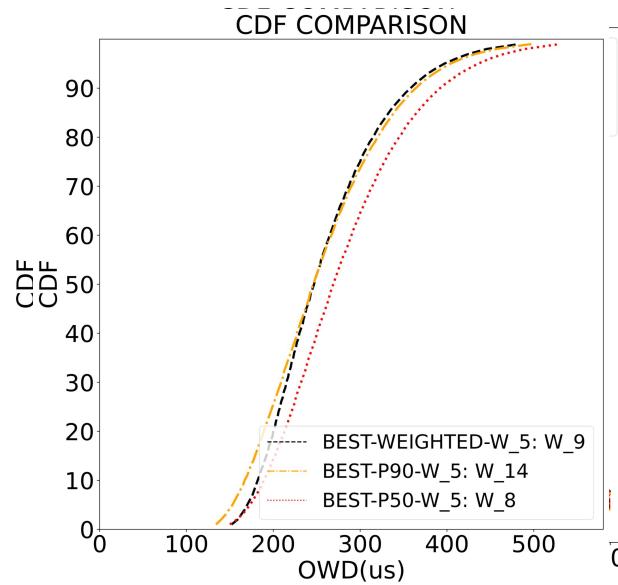


	BEST-WEIGHTED	RAND	WORST
90th %	358 μ s	383 μ s	498 μ s
median	247 μ s	256 μ s	318 μ s
mean	261.4 μ s	281.1 μ s	353.66 μ s
stddev	78.95 μ s	308.3 μ s	352.29 μ s
ID	W_9	W_15	W_16

➤ There is value in optimizing against stragglers!

Results: Strategy Comparison – WEIGHTED x P90 x P50

- **WEIGHTED** 6.53% lower 90th % than **P50**
- **P90** shows some of the best latency profiles
- **P50** shows less variance

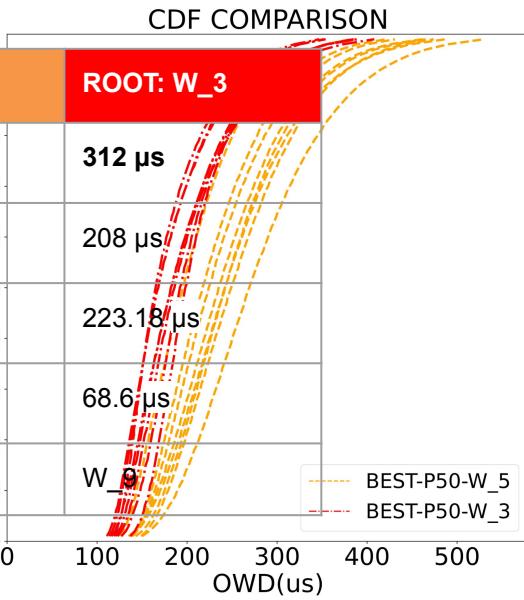
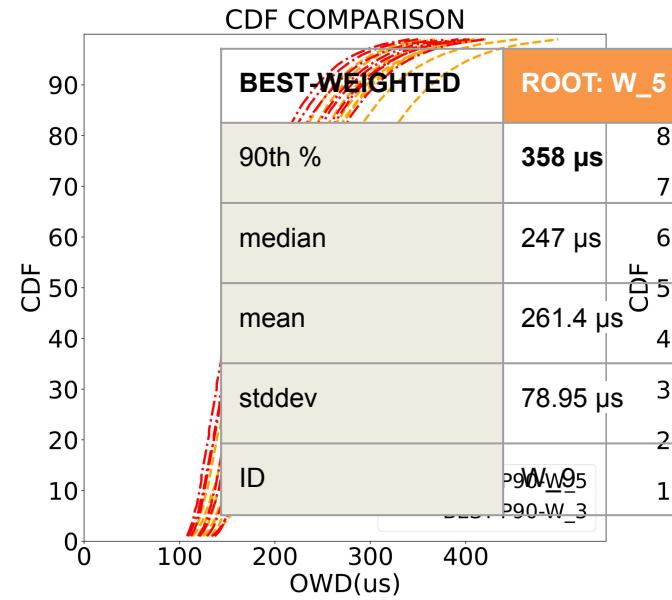
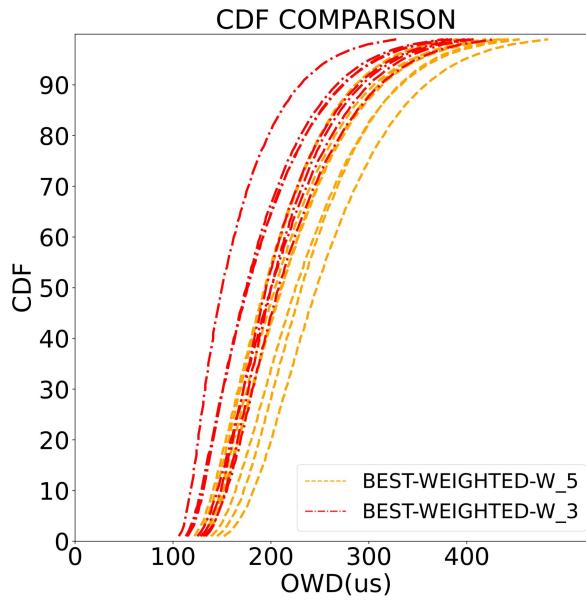


	BEST-WEIGHTED	BEST-P90	BEST-P50
90th %	358 μ s	364 μ s	392 μ s
median	247 μ s	246 μ s	269 μ s
mean	261.4 μ s	262.11 μ s	282.57 μ s
stddev	78.95 μ s	159.2 μ s	83.57 μ s
ID	W_9	W_14	W_18

➤ **P50 prioritizes stability, P90 performance, WEIGHTED seeks balance**

Results: Root Selection

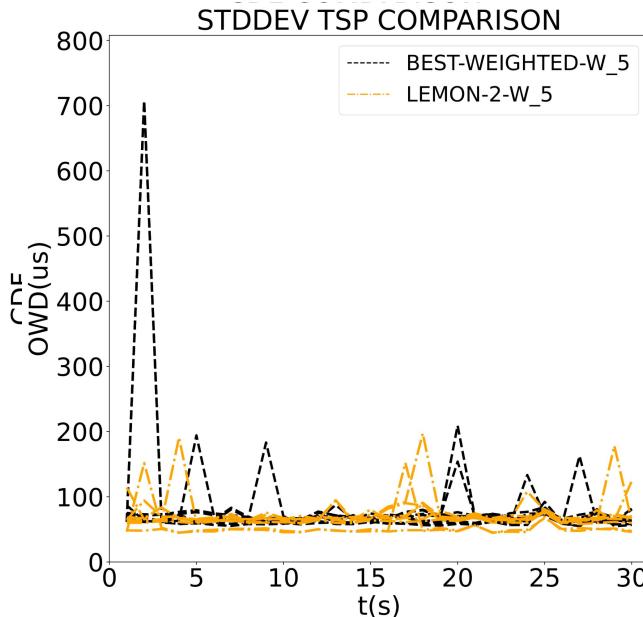
- All strategies performed worse with **W_5** as root
- **W_3**-based root shows 12.85% lower latency in 90th%



➤ Root selection impact is significant!

Results: LemonDrop Comparison

- **LEMON-2** 9.22% lower 90th % than **BEST-WEIGHTED**
- **LEMON-2** left-shifted CDF profile vs **BEST-WEIGHTED**
- **LEMON-2** slightly less variant

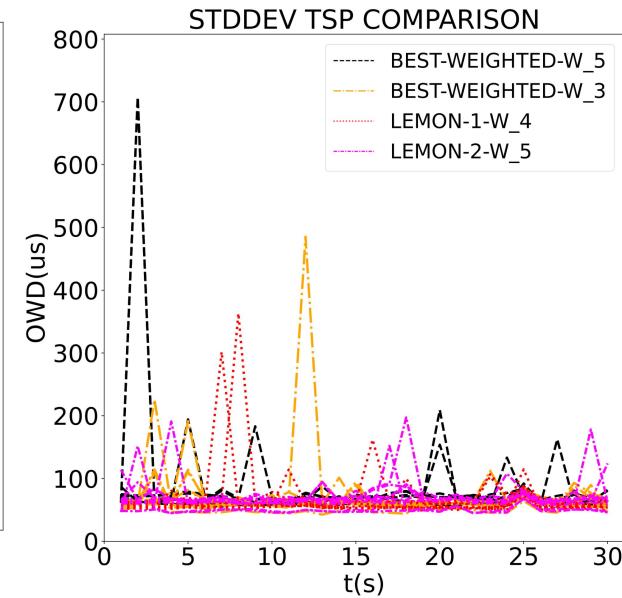
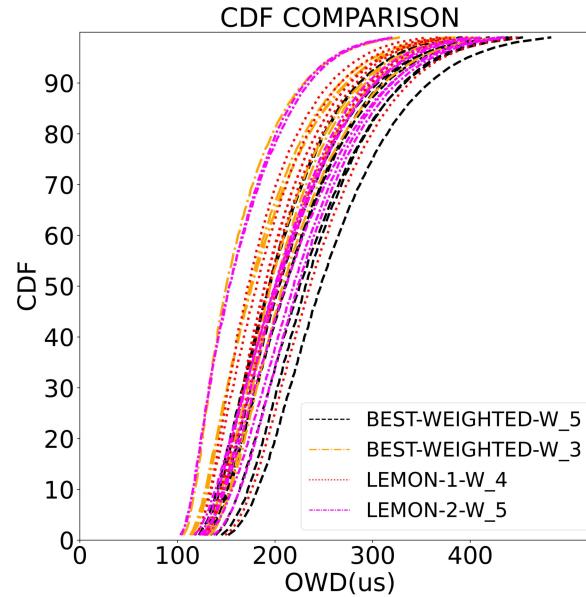
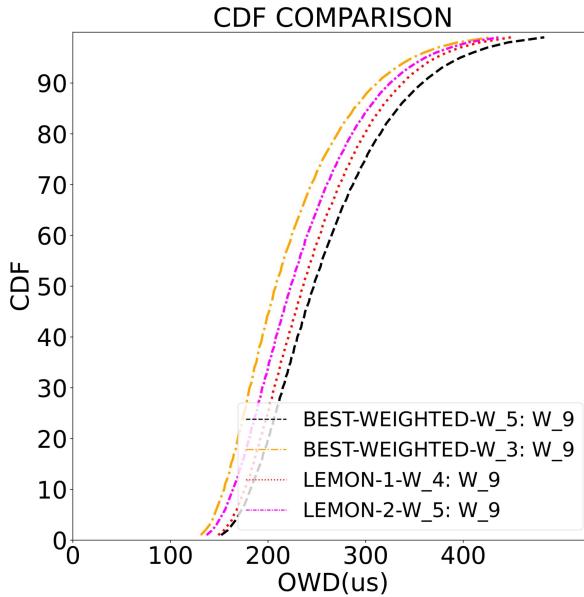


	BEST-WEIGHTED	LEMON-2
90th %	358 μs	325 μs
median	247 μs	223 μs
mean	261.4 μs	236.7 μs
stddev	78.95 μs	76.39 μs
ID	W_9	W_9

LEMON-2 converged in 0.0016sec

Results: LemonDrop Comparison

- **BEST-WEIGHTED-W_3** displays better performing worst receiver
- **LEMON-1** and **LEMON-2** show overall healthier receiver profiles and less variance



LEMON-1 and LEMON-2 converged respectively in 0.0346 and 0.0016 seconds

➤ **Comparable performance, Lemon offers more stability**

Takeaways

1. Flexible Selection Framework:
 - a. Detects stragglers
 - b. Improves multicast latency (μ s)
 - c. Trade-offs between variance x perf.
 - d. Straight-Forward, Lightweight
2. Comparable to LemonDrop
 - a. pre-defined hyperparameters (ϵ , n)

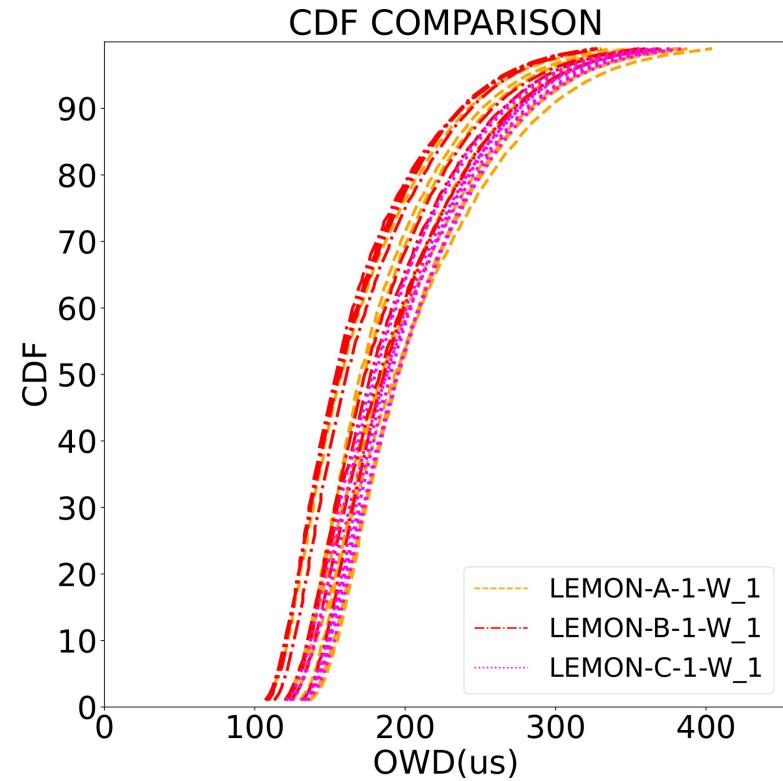
Future Work

1. Incremental Updates
 - a. Latency profile per subtree
 - b. auto-scale triggers
2. Root selection mechanism
3. Refining heuristic:
 - a. high-performance
 - b. stability

Questions!

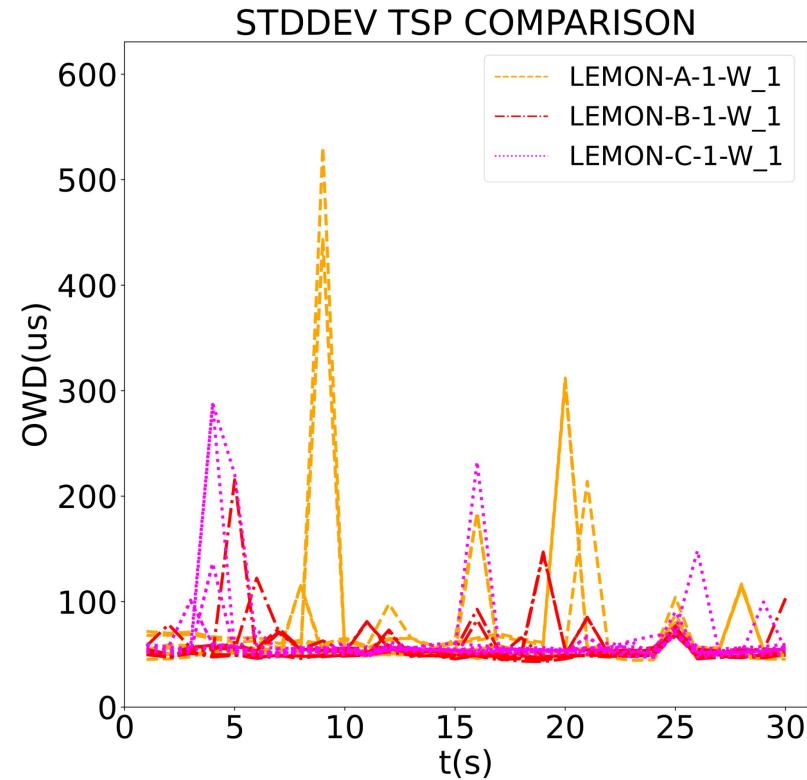
LemonDrop Parametrization

	LEMON-A	LEMON-B	LEMON-C
90th %	295 μs	273 μs	285 μs
stddev	83.35 μ s	64.19 μ s	57.64 μ s
epsilon	1×10^{-4}	2.2×10^{-5}	1.7×10^{-5}
n	1000	10 000	100 000
conv.	TRUE	FALSE	FALSE
time	0.0381 sec	11.26 sec	111.94 sec

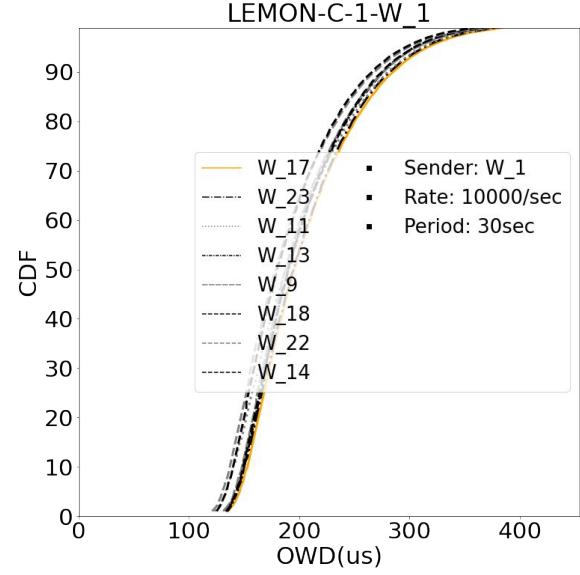
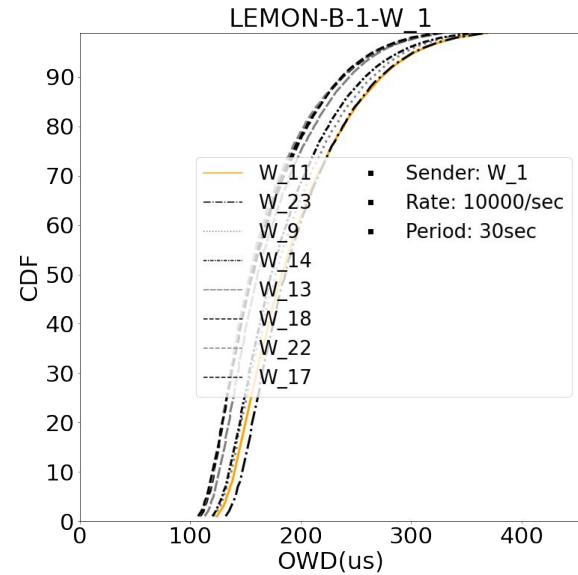
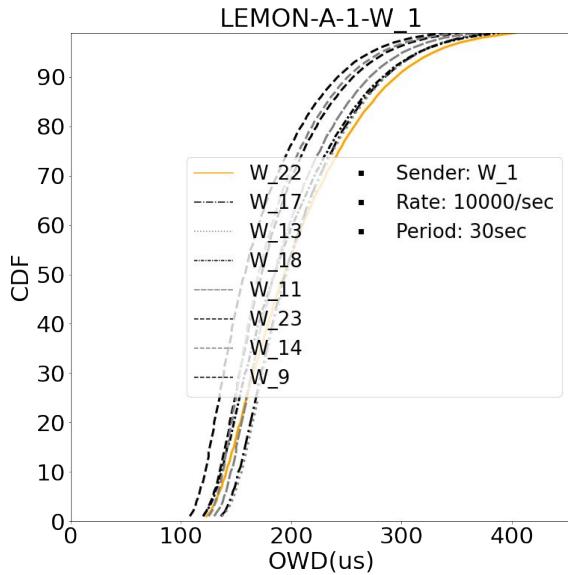


LemonDrop Parametrization

	LEMON-A	LEMON-B	LEMON-C
90th %	295 μs	273 μs	285 μs
stddev	83.35 μ s	64.19 μ s	57.64 μ s
epsilon	1×10^{-4}	2.2×10^{-5}	1.7×10^{-5}
n	1000	10 000	100 000
conv.	TRUE	FALSE	FALSE
time	0.0381 sec	11.26 sec	111.94 sec



LemonDrop Parametrization

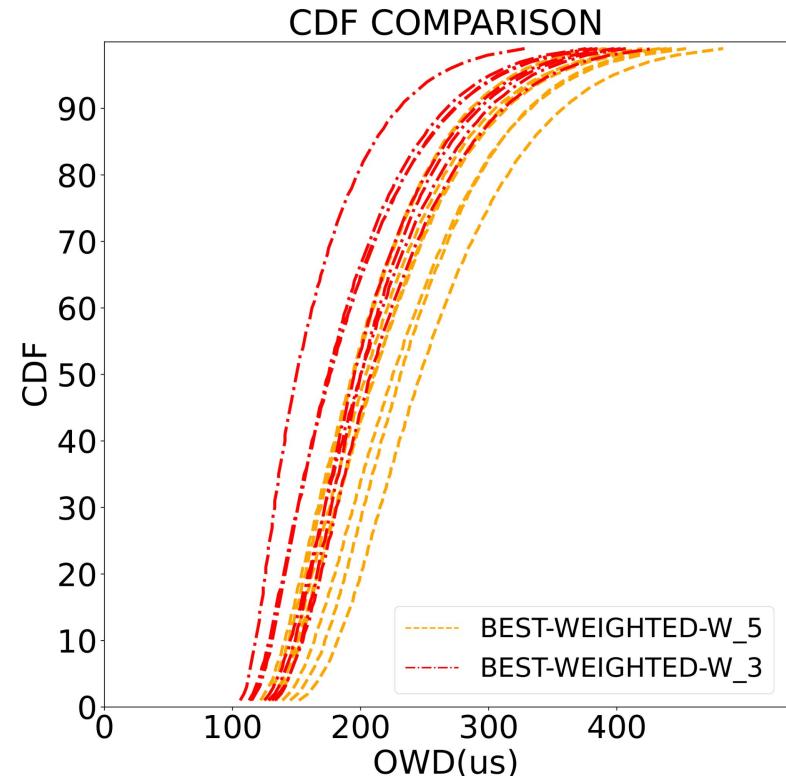


Root Selection impacts final system significantly – WEIGHTED

1. Contrast all receivers

BEST-WEIGHTED	ROOT: W_5	ROOT: W_3
90th %	358 μ s	312 μ s
median	247 μ s	208 μ s
mean	261.4 μ s	223.18 μ s
stddev	78.95 μ s	68.6 μ s
ID	W_9	W_9

$$\text{Weighted} = (0.7 \times \text{stddev}) + (0.3 \times \text{p90})$$

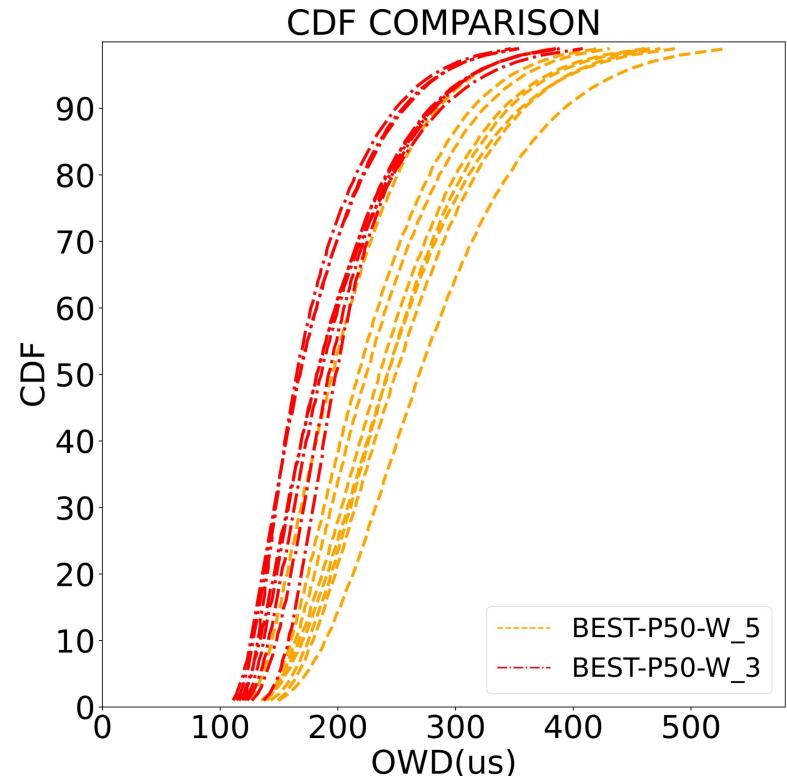


Root Selection impacts final system significantly – P50

1. Contrast all receivers

BEST-P50	ROOT: W_5	ROOT: W_3
90th %	392 μ s	288 μ s
median	269 μ s	186 μ s
mean	282.57 μ s	205.25 μ s
stddev	83.57 μ s	156.08 μ s
ID	W_18	W_9

P50: (1.0 x p50)

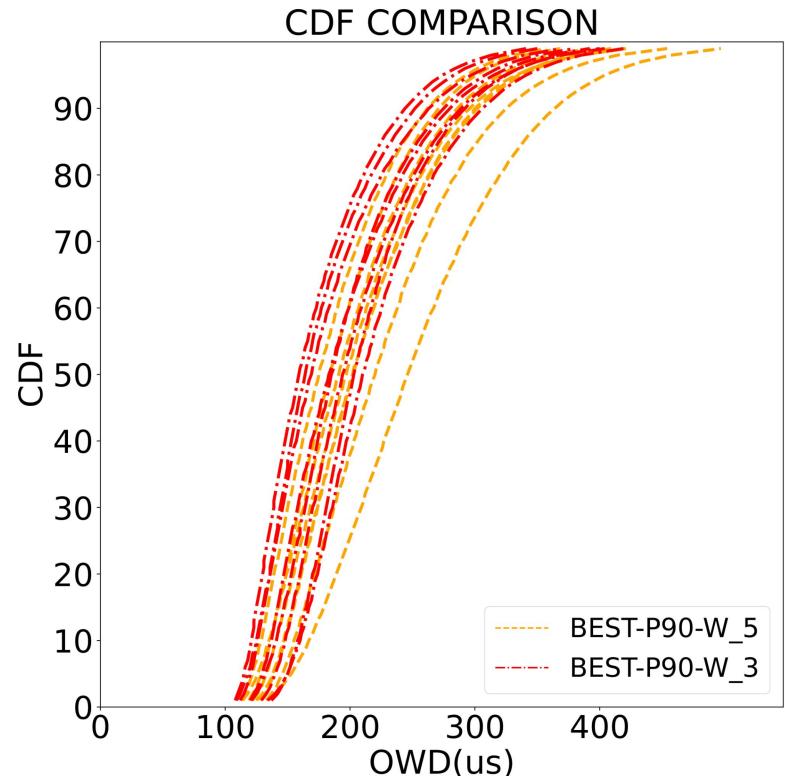


Root Selection impacts final system significantly – P90

1. Contrast all receivers

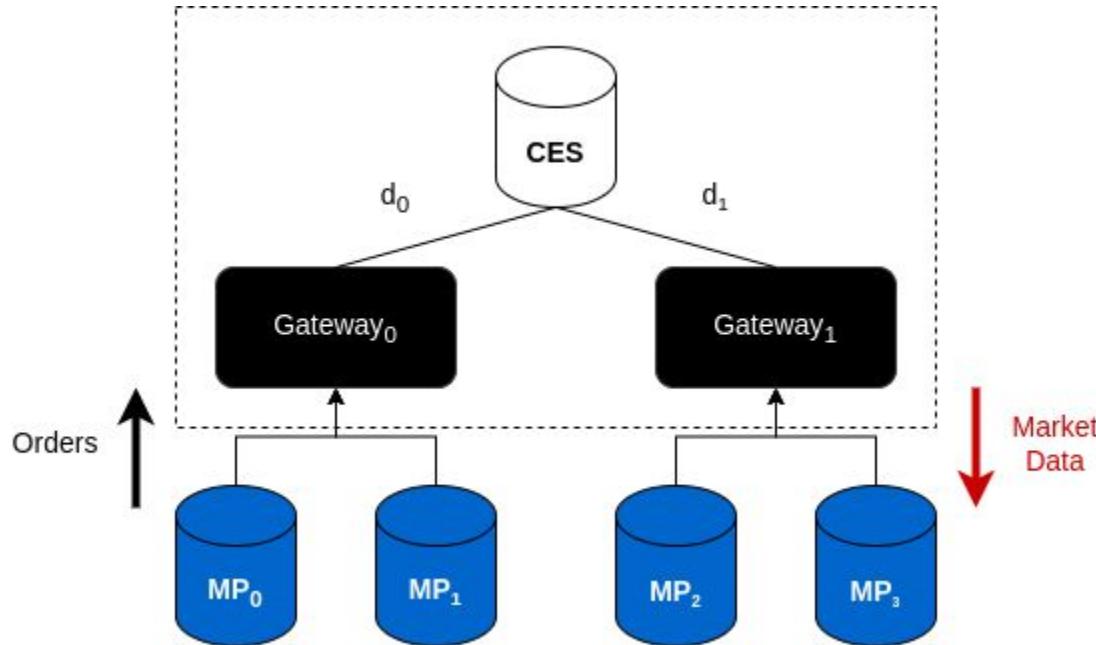
BEST-P90	ROOT: W_5	ROOT: W_3
90th %	364 μ s	306 μ s
median	246 μ s	210 μ s
mean	262.11 μ s	226.45 μ s
stddev	159.2 μ s	153.93 μ s
ID	W_14	W_6

P90: (1.0 x p90)



Background: Financial Exchanges

- Fairness or “Fair Access” to Market Participants (MPs)
- Migration to the Cloud: i) **Mechanisms for fair delivery** ii) **Latency Variance**



Background: Financial Exchanges to the Cloud

BIDS TO BUY

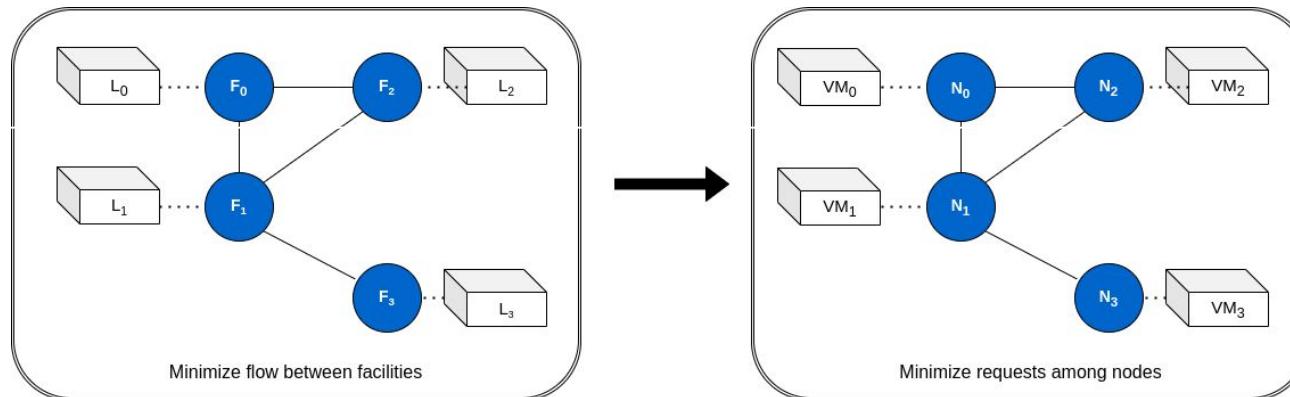
\$29.49	100	100	200
\$29.48	50	30	
\$29.47	100		
\$29.46	50	100	100
\$29.45	200		

OFFERS TO SELL

\$29.54	100	200
\$29.53	50	
\$29.52	40	50
\$29.51	50	50
\$29.50	100	100

State-of-the-Art: LemonDrop¹

- Core of method:
 - Clock Synchronization
 - (QAP) Optimization Method: Facilities=Nodes, VMs=Locations
- Λ : Load Matrix, Δ : OWD Matrix, P : Permutation Matrix
- $D = \sum_{i,j=1}^L \lambda_{ij} d_{ij}$, $D(P) = \text{trace}(\Lambda P \Delta^T P^T)$



¹LemonDrop: <https://searchworks.stanford.edu/view/14423035>

State-of-the-Art: LemonDrop¹

- Λ : Load Matrix, Δ : OWD Matrix, P : Permutation Matrix
- Hadamard Product: $D = \sum_{i,j=1}^L \lambda_{ij} d_{ij}$
- Goal: Minimize $D(P) = \text{trace}(\Lambda P \Delta^T P^T)$

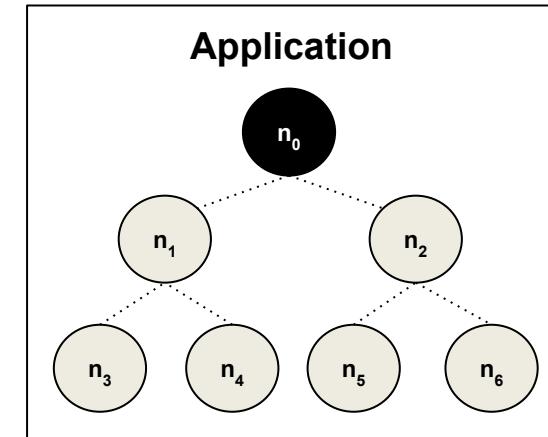
$$\begin{array}{c} \Delta: (N \times N) \\ \left[\begin{array}{ccc} d_{00} & d_{01} & d_{02} \\ d_{10} & d_{11} & d_{12} \\ d_{20} & d_{21} & d_{22} \end{array} \right] \end{array} \quad \begin{array}{c} \Lambda: (L \times L) \\ \left[\begin{array}{ccc} \lambda_{00} & \lambda_{01} & \lambda_{02} \\ \lambda_{10} & \lambda_{11} & \lambda_{12} \end{array} \right] \end{array} \quad \xrightarrow{\text{trace}} \quad \begin{array}{c} P: (N \times N) \\ \left(\Lambda \left[\begin{array}{ccc} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{array} \right] \Delta^T P^T \right) \end{array}$$

¹LemonDrop: <https://searchworks.stanford.edu/view/14423035>

State-of-the-Art: LemonDrop¹

- Λ : Load Matrix, Δ : OWD Matrix, P : Permutation Matrix
- Hadamard Product: $D = \sum_{i,j=1}^L \lambda_{ij} d_{ij}$
- Goal: Minimize $D(P) = \text{trace}(\Lambda P \Delta^T P^T)$

$$\begin{array}{c} \Delta: (10 \times 10) \\ \left[\begin{array}{ccccccc} d_{00} & d_{01} & \dots & d_{09} \\ d_{10} & d_{11} & \dots & d_{19} \\ \dots \\ d_{90} & d_{91} & \dots & d_{99} \end{array} \right] \end{array} \quad \begin{array}{c} \Lambda: (7 \times 7) \\ \left[\begin{array}{ccccccc} 0 & 1 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 1 \\ \dots \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{array} \right] \end{array} \quad \longleftrightarrow$$



¹LemonDrop: <https://searchworks.stanford.edu/view/14423035>

State-of-the-Art: LemonDrop¹

- Λ : Load Matrix, Δ : OWD Matrix, P : Permutation Matrix
- $D = \sum_{i,j=1}^L \lambda_{ij} d_{ij}$, $D(P) = \text{trace}(\Lambda P \Delta^T P^T)$

Algorithm 1 FAQ Algorithm [136]

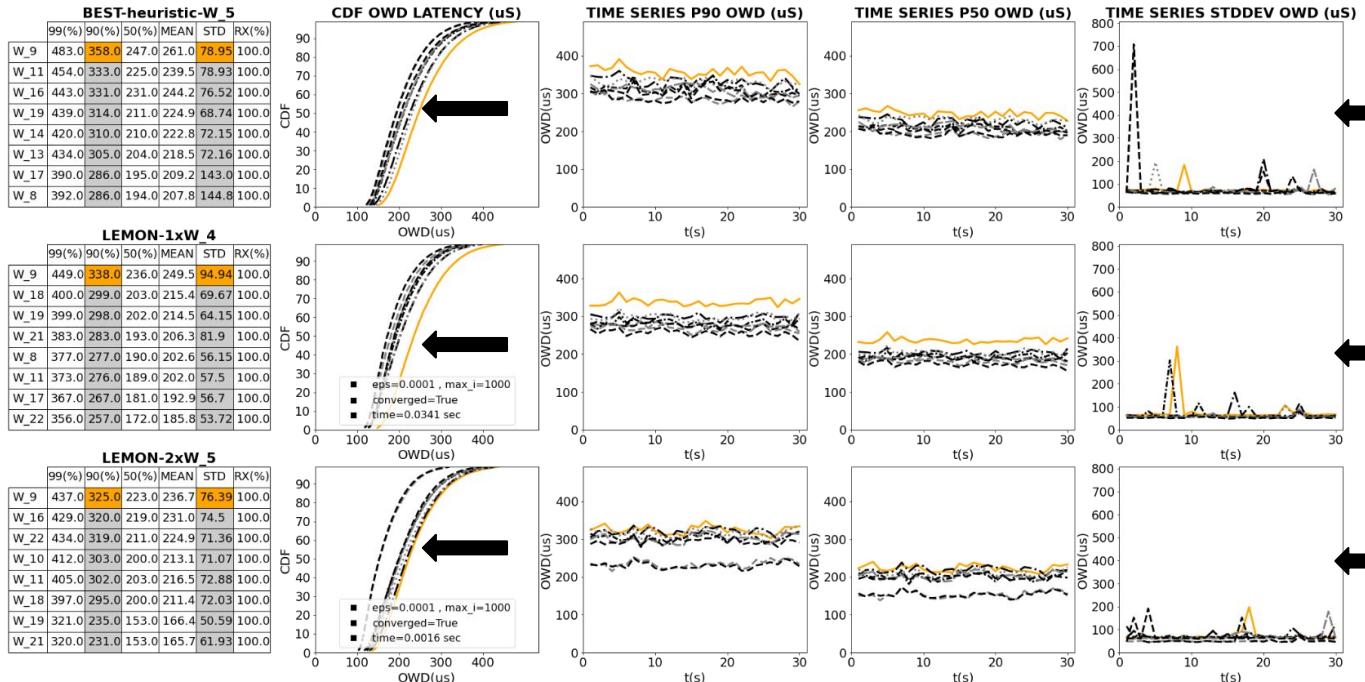
```
function FAQ( $\Lambda, \Delta, \epsilon, n$ )
     $P^{(0)} = 11^T$ 
     $i = 0$                                  $\triangleright$  Iteration number
     $s = 0$                                  $\triangleright$  Stopping condition
    while  $s = 0$  do
         $\nabla f(P^{(i)}) = -\Lambda P^{(i)} \Delta^T - \Lambda^T P^{(i)} \Delta$            $\triangleright$  Gradient wrt current solution
         $Q^{(i)} = \operatorname{argmin}_{P \in \mathcal{D}} \text{trace}(\nabla f(P^{(i)})^T P)$        $\triangleright$  Direction which minimizes 1st order approx of  $f(P)$ 
         $\alpha^i = \min_{\alpha \in [0,1]} f(P^{(i)} + \alpha Q^{(i)})$                        $\triangleright$  Best step size in chosen direction
         $P^{(i+1)} = P^{(i)} + \alpha^i Q^{(i)}$                                       $\triangleright$  Update our current solution
        if  $\|P^{(i)} - P^{(i-1)}\|_F < \epsilon$  then  $s = 1$                           $\triangleright$  Stop opt
        if  $i > n$  then  $s = 1$                                           $\triangleright$  Stop opt
         $i = i + 1$ 
     $P^{(\text{final})} = \operatorname{argmin}_{P \in \mathcal{P}} P^{(i-1)} P^T$ 
    return  $P^{(\text{final})}$ 
```

¹LemonDrop: <https://searchworks.stanford.edu/view/14423035>

Results

EVALUATION COMPARISON - BEST x LEMONS

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=30 sec



- LemonDrop(Epsilon, MAX_I): Smaller Epsilon => Pickier Solution
- HEURISTIC shows comparable performance to LemonDrop => 358us vs 338us vs 325us

Background: Jasper

10 Receivers		
D	F	OML (μ s)
1	10	66
2	4	88

100 Receivers		
D	F	OML (μ s)
1	100	351
2	10	139
3	5	141

1000 Receivers		
D	F	OML (μ s)
2	32	282
3	10	217
4	6	226

- $D = \log_{10}(N)$
- $N = 100, D=2, F=10$
- $N = 1000, D=3, F=10$

Results

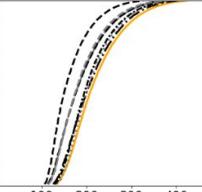
EVALUATION COMPARISON - BEST x LEMONS

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=30 sec

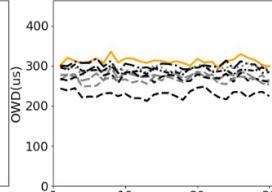
BEST-heuristic-W_3

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_9	428.0	312.0	208.0	223.1	68.6	100.0
W_4	407.0	300.0	200.0	213.8	67.13	100.0
W_19	403.0	293.0	200.0	214.1	62.75	100.0
W_1	398.0	289.0	195.0	209.6	72.58	100.0
W_10	385.0	274.0	176.0	191.6	70.17	100.0
W_17	381.0	273.0	176.0	192.2	109.5	100.0
W_8	376.0	267.0	175.0	189.4	64.88	100.0
W_13	328.0	230.0	150.0	164.6	56.76	100.0

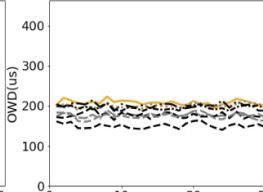
CDF OWD LATENCY (uS)



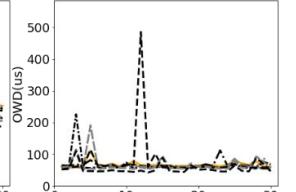
TIME SERIES P90 OWD (uS)



TIME SERIES P50 OWD (uS)

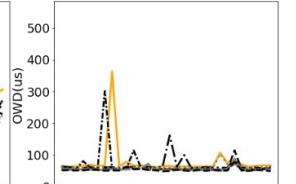
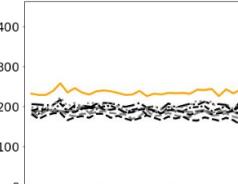
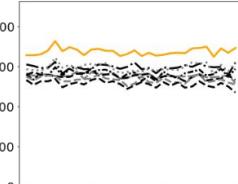
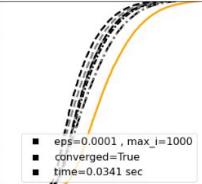


TIME SERIES STDDEV OWD (uS)



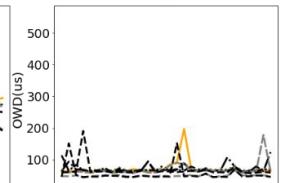
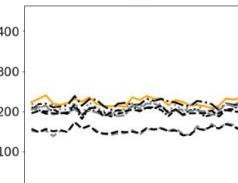
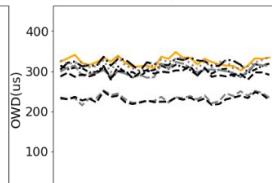
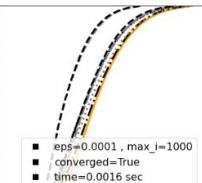
LEMON-1xW 4

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_9	449.0	338.0	236.0	249.5	94.94	100.0
W_18	400.0	299.0	203.0	215.4	69.67	100.0
W_19	399.0	298.0	202.0	214.5	64.15	100.0
W_21	383.0	283.0	193.0	206.3	81.9	100.0
W_8	377.0	277.0	190.0	202.6	56.19	100.0
W_11	373.0	276.0	189.0	202.0	57.5	100.0
W_17	367.0	267.0	181.0	192.9	56.7	100.0
W_22	356.0	257.0	172.0	185.8	53.72	100.0



LEMON-2xW 5

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_9	437.0	325.0	223.0	236.7	76.39	100.0
W_16	429.0	320.0	219.0	231.0	74.5	100.0
W_22	434.0	319.0	211.0	224.9	71.36	100.0
W_10	412.0	303.0	200.0	213.1	71.07	100.0
W_11	405.0	302.0	203.0	216.5	72.88	100.0
W_18	397.0	295.0	200.0	211.4	72.03	100.0
W_19	321.0	235.0	153.0	166.4	50.59	100.0
W_21	320.0	231.0	153.0	165.7	61.93	100.0



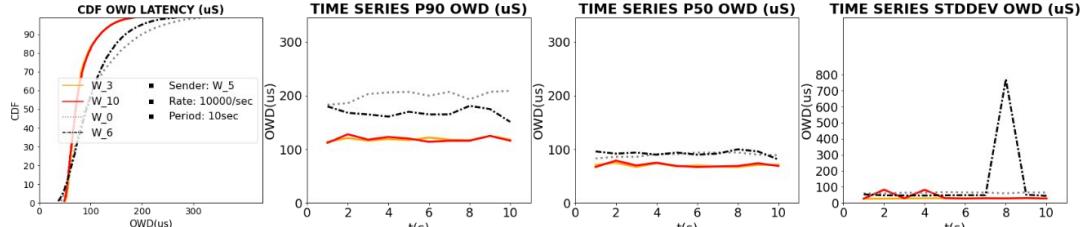
Results

BEST HEURISTIC: STAGES 1-3

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=10 sec

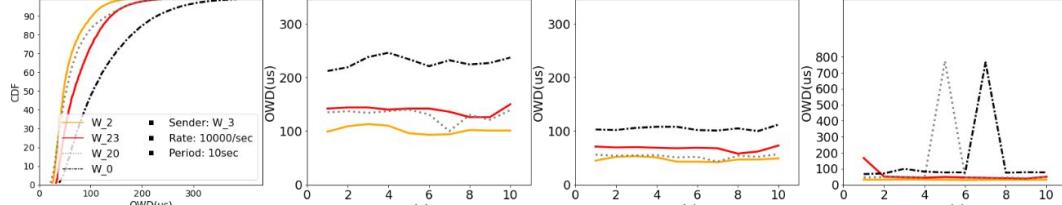
BEST-heuristic-W_5 STAGE 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_3	181.0	119.0	71.0	80.02	28.4	100.0
W_10	186.0	119.0	71.0	80.76	45.22	100.0
W_0	322.0	200.0	90.0	111.3	63.95	100.0
W_6	268.0	168.0	92.0	110.4	249.4	100.0



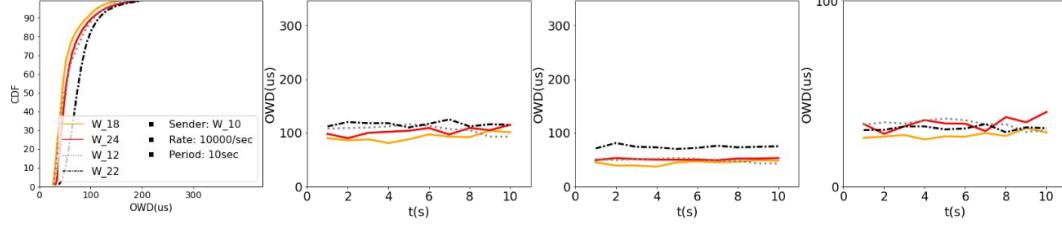
BEST-heuristic-W_5 STAGE 2

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_2	169.0	103.0	48.0	58.51	31.59	100.0
W_23	237.0	140.0	68.0	82.16	69.35	100.0
W_20	237.0	132.0	53.0	75.19	249.3	100.0
W_0	384.0	229.0	105.0	133.4	254.1	100.0



BEST-heuristic-W_5 STAGE 3

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_18	154.0	93.0	45.0	53.8	27.98	100.0
W_24	192.0	103.0	51.0	62.39	34.24	100.0
W_12	180.0	107.0	49.0	61.07	33.94	100.0
W_22	199.0	117.0	74.0	80.54	31.53	100.0



- Simple quick selection seems to at least prune bad VMs

Results

BEST HEURISTIC: STAGES 4-7

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=10 sec

BEST-heuristic-W_5 STAGE 4

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_17	201.0	116.0	58.0	69.77	36.24	100.0
W_8	198.0	112.0	56.0	67.73	41.63	100.0
W_9	239.0	134.0	82.0	90.96	37.88	100.0
W_13	153.0	99.0	62.0	69.52	37.59	100.0

BEST-heuristic-W_5 STAGE 5

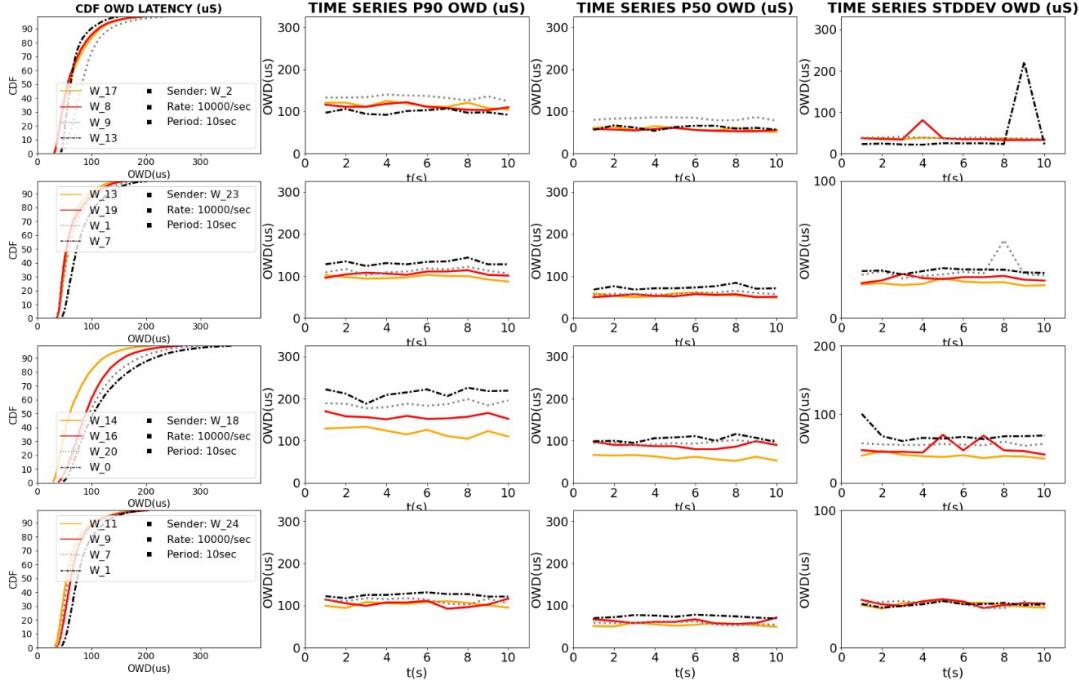
	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_13	151.0	97.0	55.0	63.11	25.83	100.0
W_19	162.0	106.0	54.0	64.45	29.1	100.0
W_1	187.0	113.0	58.0	69.61	35.41	100.0
W_7	202.0	132.0	73.0	83.96	34.76	100.0

BEST-heuristic-W_5 STAGE 6

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_14	211.0	122.0	60.0	72.08	39.73	100.0
W_16	265.0	158.0	89.0	100.1	51.49	100.0
W_20	300.0	187.0	96.0	110.5	56.43	100.0
W_0	360.0	214.0	103.0	124.3	70.57	100.0

BEST-heuristic-W_5 STAGE 7

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_11	179.0	103.0	54.0	63.96	31.53	100.0
W_9	204.0	105.0	62.0	70.88	32.73	100.0
W_7	181.0	112.0	58.0	68.75	32.05	100.0
W_1	194.0	125.0	74.0	82.79	31.6	100.0

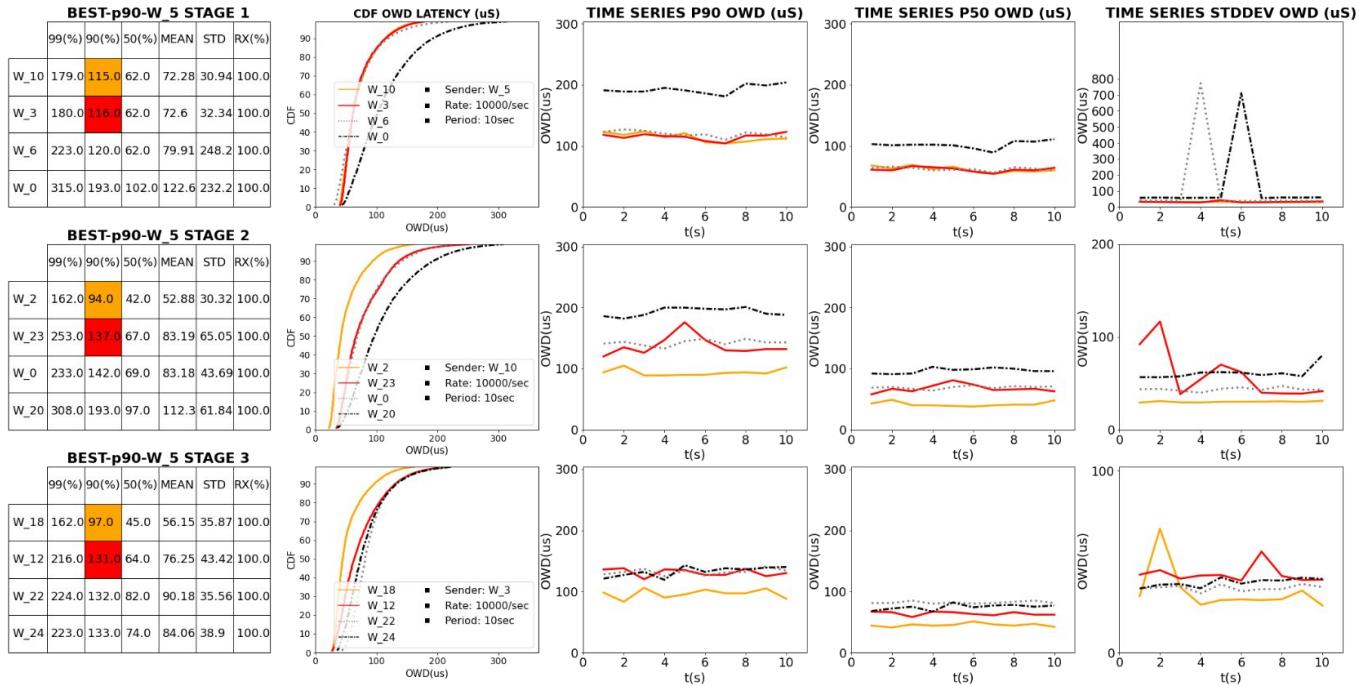


- Simple quick selection seems to at least prune bad VMs

Results

BEST p90: STAGES 1-3

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=10 sec

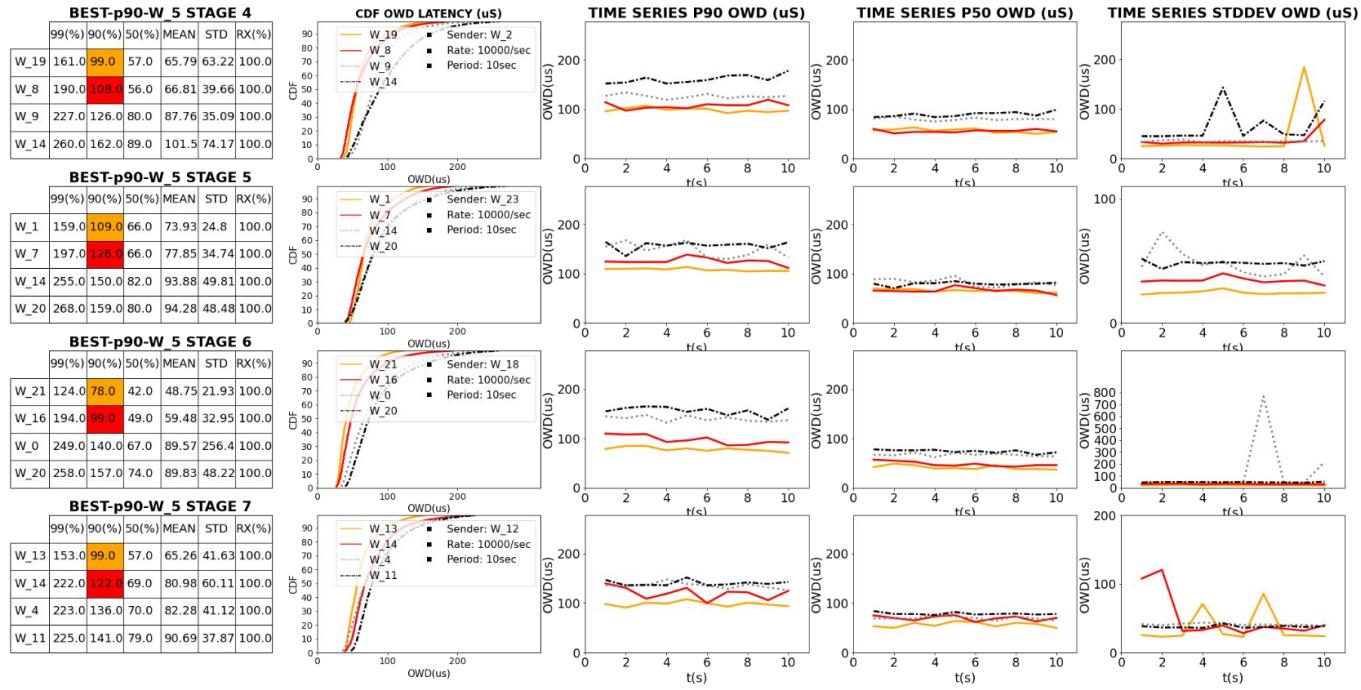


- Simple quick selection seems to at least prune bad VMs

Results

BEST p90: STAGES 4-7

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=10 sec

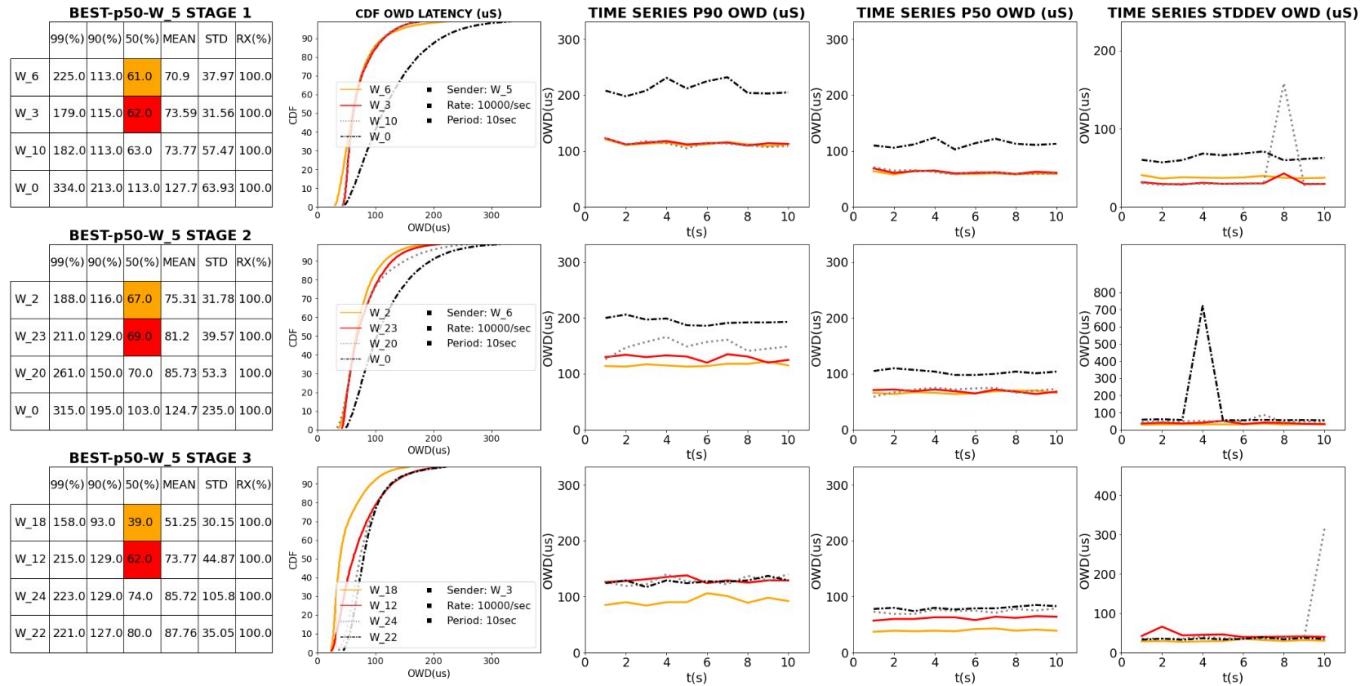


- Simple quick selection seems to at least prune bad VMs

Results

BEST p50: STAGES 1-3

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=10 sec



- Simple quick selection seems to at least prune bad VMs

Results

BEST p50: STAGES 4-7

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=10 sec

BEST-p50-W_5 STAGE 4

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_19	145.0	85.0	46.0	53.78	24.37	100.0
W_10	164.0	99.0	57.0	65.83	26.05	100.0
W_9	204.0	97.0	57.0	65.25	32.23	100.0
W_14	229.0	132.0	66.0	83.09	146.8	100.0

BEST-p50-W_5 STAGE 5

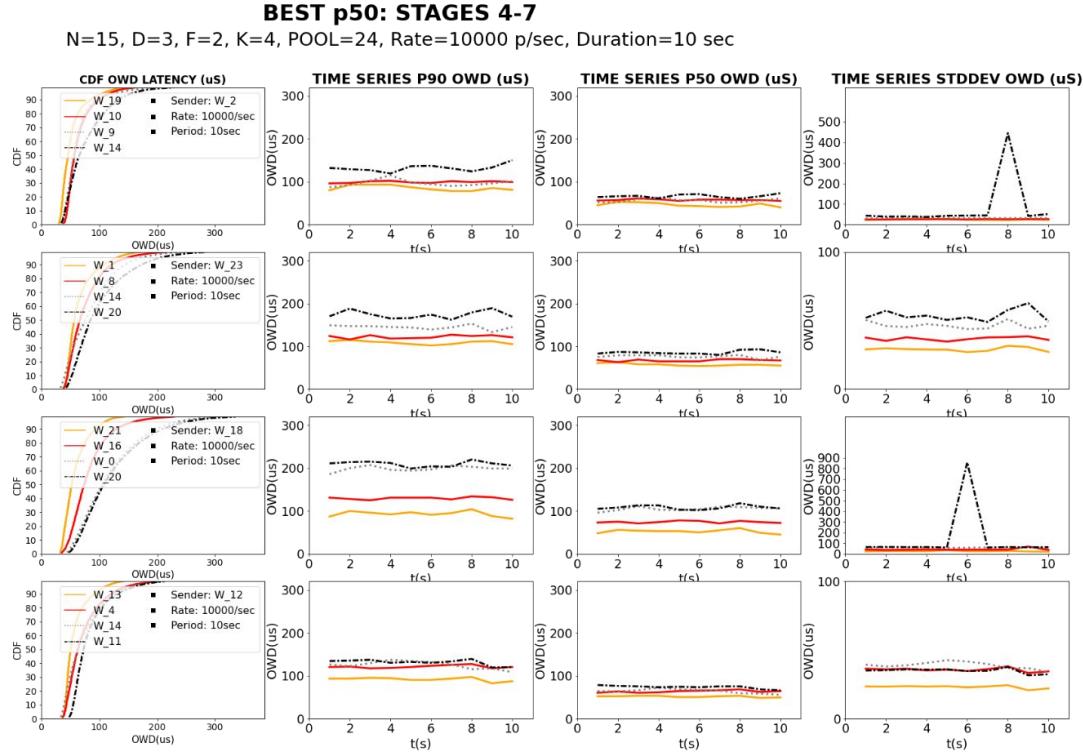
	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_1	171.0	109.0	57.0	67.75	29.2	100.0
W_8	215.0	122.0	67.0	77.82	36.9	100.0
W_14	244.0	145.0	76.0	86.82	46.65	100.0
W_20	283.0	174.0	85.0	102.0	54.03	100.0

BEST-p50-W_5 STAGE 6

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_21	142.0	93.0	53.0	60.17	25.74	100.0
W_16	229.0	129.0	74.0	83.87	42.4	100.0
W_0	310.0	199.0	105.0	119.5	59.12	100.0
W_20	335.0	210.0	108.0	133.1	278.5	100.0

BEST-p50-W_5 STAGE 7

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_13	141.0	92.0	51.0	59.58	23.11	100.0
W_4	204.0	121.0	63.0	74.72	35.6	100.0
W_14	216.0	125.0	64.0	75.6	39.18	100.0
W_11	211.0	131.0	73.0	84.8	35.03	100.0



- Simple quick selection seems to at least prune bad VMs

Results

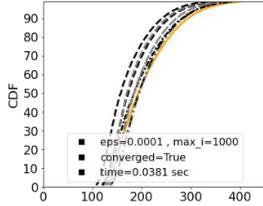
LEMON PARAMETRIZATION

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=30 sec

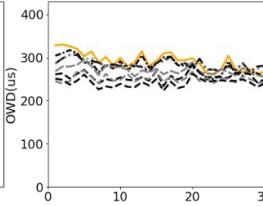
LEMON-A-1xW 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_22	404.0	295.0	193.0	209.1	88.35	100.0
W_17	382.0	285.0	195.0	209.8	99.75	100.0
W_13	380.0	285.0	196.0	212.1	99.95	100.0
W_18	389.0	283.0	185.0	200.3	85.03	100.0
W_11	368.0	271.0	185.0	199.4	76.68	100.0
W_23	353.0	256.0	173.0	186.1	52.02	100.0
W_14	349.0	251.0	170.0	185.2	62.71	100.0
W_9	338.0	242.0	158.0	173.2	109.6	100.0

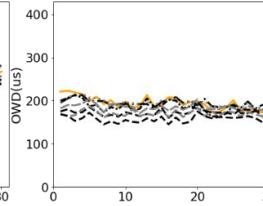
CDF OWD LATENCY (uS)



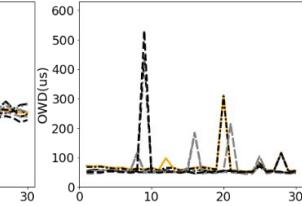
TIME SERIES P90 OWD (uS)



TIME SERIES P50 OWD (uS)

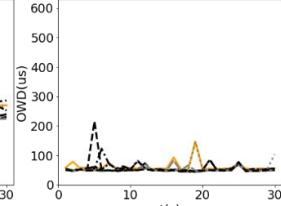
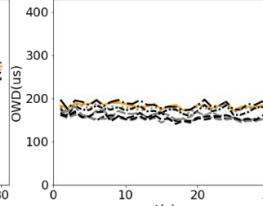
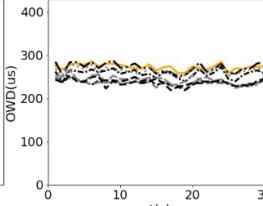
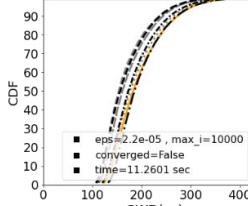


TIME SERIES STDDEV OWD (uS)



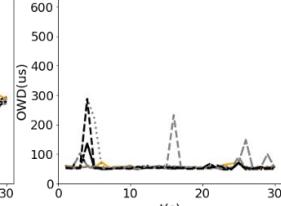
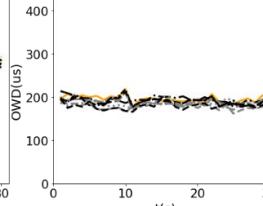
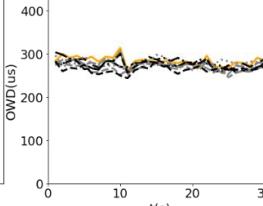
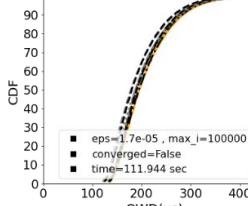
LEMON-B-1xW 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_11	369.0	273.0	183.0	197.1	64.19	100.0
W_23	369.0	272.0	186.0	200.0	60.13	100.0
W_9	360.0	265.0	177.0	191.2	62.17	100.0
W_14	356.0	258.0	173.0	187.3	56.36	100.0
W_13	330.0	244.0	161.0	175.2	51.33	100.0
W_18	328.0	238.0	155.0	169.8	63.81	100.0
W_22	327.0	237.0	154.0	168.0	51.5	100.0
W_17	325.0	237.0	157.0	169.6	50.21	100.0



LEMON-C-1xW 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_17	384.0	285.0	195.0	209.3	57.64	100.0
W_23	378.0	282.0	195.0	207.4	57.87	99.88
W_11	376.0	280.0	188.0	202.0	86.96	99.88
W_13	369.0	276.0	190.0	203.9	58.18	100.0
W_9	372.0	276.0	188.0	202.1	62.44	100.0
W_18	367.0	275.0	190.0	203.1	58.47	100.0
W_22	364.0	268.0	179.0	193.1	71.42	99.88
W_14	359.0	266.0	180.0	194.0	74.64	99.88



- Smaller epsilon: negligible performance enhancement & too long to converge

Results

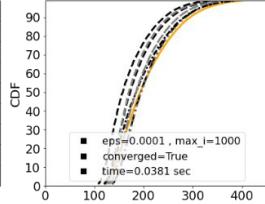
LEMON-A RUNS

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=30 sec

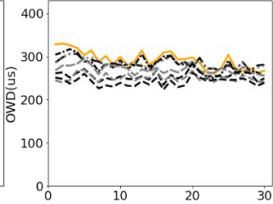
LEMON-A-1xW 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_22	404.0	295.0	193.0	209.1	88.35	100.0
W_17	382.0	285.0	195.0	209.8	99.75	100.0
W_13	380.0	285.0	196.0	212.1	99.95	100.0
W_18	389.0	283.0	185.0	200.3	85.03	100.0
W_11	368.0	271.0	185.0	199.4	76.68	100.0
W_23	353.0	256.0	173.0	186.1	52.02	100.0
W_14	349.0	251.0	170.0	185.2	62.71	100.0
W_9	338.0	242.0	158.0	173.2	109.6	100.0

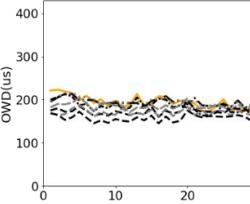
CDF OWD LATENCY (uS)



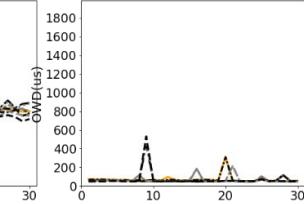
TIME SERIES P90 OWD (uS)



TIME SERIES P50 OWD (uS)

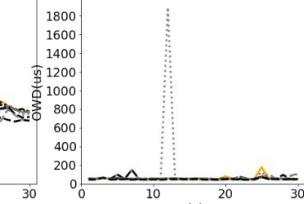
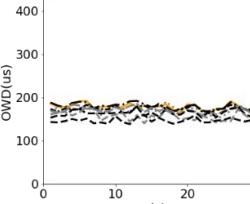
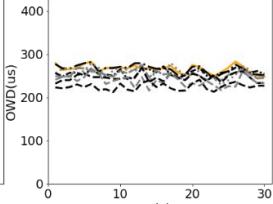
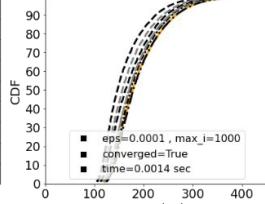


TIME SERIES STDDEV OWD (uS)



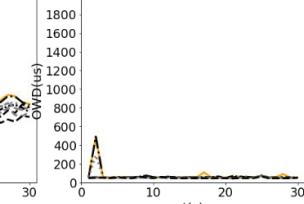
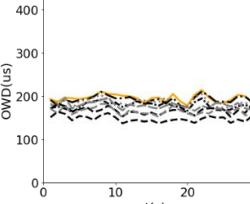
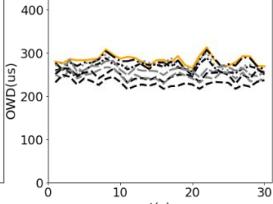
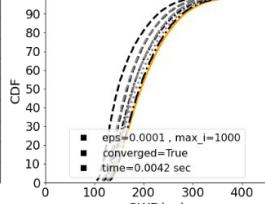
LEMON-A-2xW 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_15	364.0	267.0	178.0	192.7	63.25	100.0
W_18	363.0	266.0	179.0	193.3	54.49	100.0
W_10	358.0	259.0	171.0	194.1	35.31	100.0
W_13	348.0	258.0	174.0	187.9	51.97	100.0
W_22	345.0	254.0	171.0	184.4	60.8	100.0
W_17	338.0	247.0	163.0	177.4	62.83	100.0
W_23	334.0	240.0	157.0	171.0	59.29	100.0
W_14	314.0	225.0	147.0	160.7	48.19	100.0



LEMON-A-3xW 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_23	385.0	284.0	195.0	209.3	108.2	100.0
W_14	378.0	279.0	191.0	206.0	106.0	100.0
W_15	370.0	274.0	185.0	199.0	76.79	100.0
W_22	363.0	267.0	178.0	191.7	56.04	100.0
W_13	353.0	262.0	179.0	192.6	51.79	100.0
W_17	338.0	249.0	167.0	180.6	50.78	100.0
W_18	338.0	247.0	165.0	178.9	51.18	100.0
W_9	318.0	230.0	148.0	162.8	49.84	100.0



Results

LEMON-B RUNS

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=30 sec

LEMON-B-1xW 1

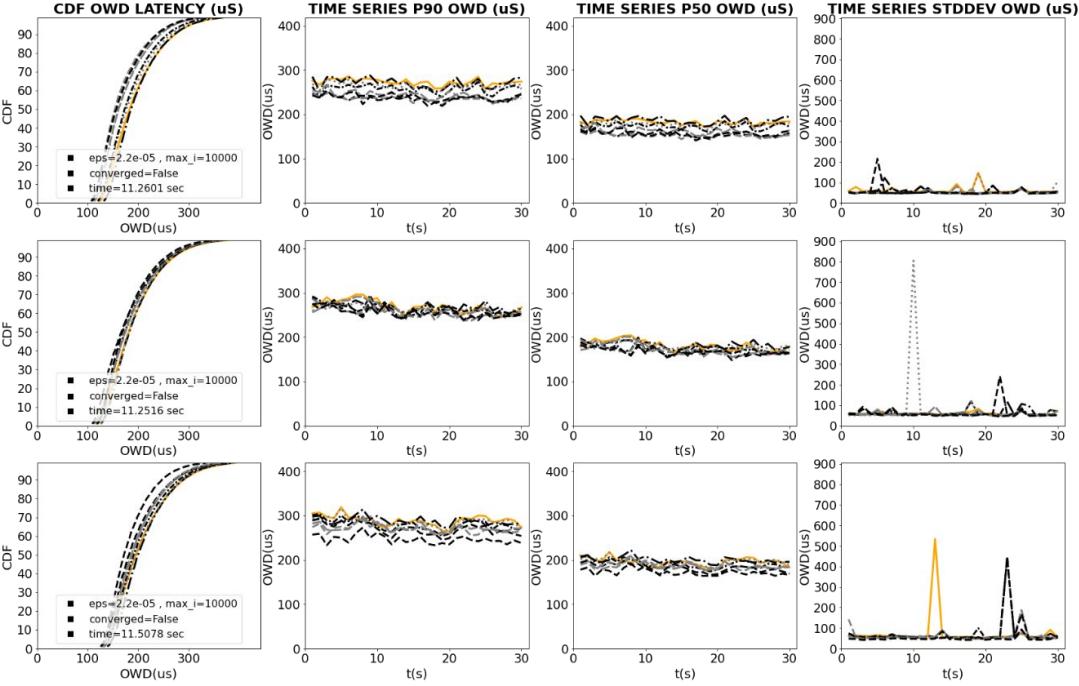
	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_11	369.0	273.0	183.0	197.1	64.19	100.0
W_23	369.0	272.0	186.0	200.0	60.13	100.0
W_9	360.0	265.0	177.0	191.2	62.17	100.0
W_14	356.0	258.0	173.0	187.3	56.36	100.0
W_13	330.0	244.0	161.0	175.2	51.33	100.0
W_18	328.0	238.0	155.0	169.8	63.81	100.0
W_22	327.0	237.0	154.0	168.0	51.5	100.0
W_17	325.0	237.0	157.0	169.6	50.21	100.0

LEMON-B-2xW 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_18	377.0	271.0	182.0	194.9	61.16	100.0
W_13	372.0	269.0	182.0	196.8	66.96	100.0
W_9	368.0	266.0	181.0	197.7	159.2	100.0
W_22	366.0	265.0	174.0	188.4	57.38	100.0
W_17	371.0	264.0	175.0	188.3	61.7	100.0
W_15	360.0	261.0	167.0	182.0	74.44	100.0
W_23	355.0	257.0	169.0	184.1	53.54	100.0
W_14	349.0	255.0	168.0	183.1	55.6	100.0

LEMON-B-3xW 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_22	392.0	290.0	195.0	211.2	114.5	100.0
W_13	385.0	289.0	200.0	213.4	58.22	100.0
W_23	387.0	287.0	196.0	209.3	57.74	100.0
W_17	384.0	281.0	191.0	206.7	100.7	100.0
W_18	370.0	273.0	186.0	200.0	65.45	100.0
W_9	370.0	271.0	182.0	197.5	102.1	100.0
W_15	360.0	270.0	187.0	200.3	59.33	100.0
W_14	334.0	247.0	172.0	184.8	51.02	100.0



Results

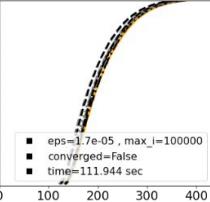
LEMON-C RUNS

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=30 sec

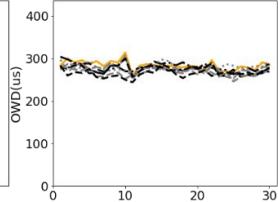
LEMON-C-1xW 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_17	384.0	285.0	195.0	209.3	57.64	100.0
W_23	378.0	282.0	195.0	207.4	57.87	99.88
W_11	376.0	280.0	188.0	202.0	86.96	99.88
W_13	369.0	276.0	190.0	203.9	58.18	100.0
W_9	372.0	276.0	188.0	202.1	62.44	100.0
W_18	367.0	275.0	190.0	203.1	58.47	100.0
W_22	364.0	268.0	179.0	193.1	71.42	99.88
W_14	359.0	266.0	180.0	194.0	74.64	99.88

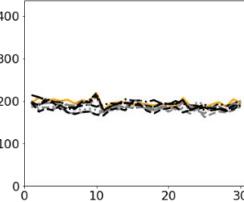
CDF OWD LATENCY (uS)



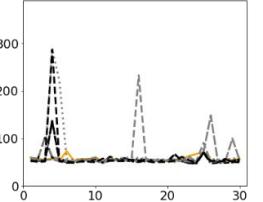
TIME SERIES P90 OWD (uS)



TIME SERIES P50 OWD (uS)



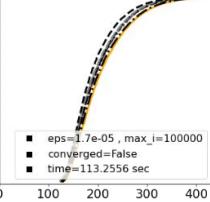
TIME SERIES STDDEV OWD (uS)



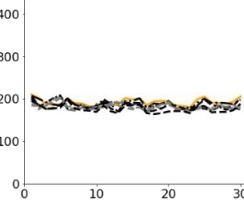
LEMON-C-2xW 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_17	399.0	287.0	192.0	207.8	72.61	100.0
W_13	394.0	283.0	189.0	204.8	73.49	100.0
W_9	370.0	273.0	184.0	199.1	75.69	100.0
W_18	366.0	271.0	184.0	197.4	57.55	100.0
W_15	365.0	270.0	182.0	196.2	72.74	100.0
W_23	377.0	270.0	184.0	198.6	71.22	100.0
W_11	364.0	268.0	181.0	194.8	70.82	100.0
W_14	364.0	257.0	176.0	190.7	69.63	100.0

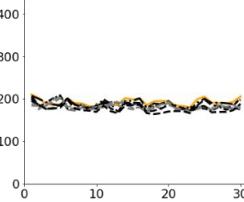
CDF OWD LATENCY (uS)



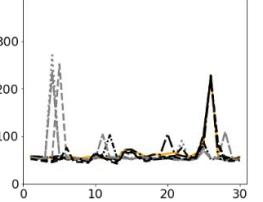
TIME SERIES P90 OWD (uS)



TIME SERIES P50 OWD (uS)



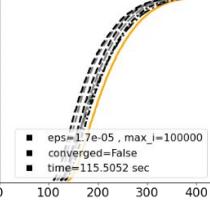
TIME SERIES STDDEV OWD (uS)



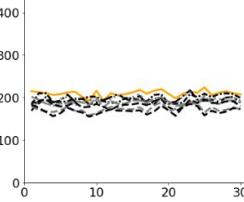
LEMON-C-3xW 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_13	403.0	303.0	210.0	223.0	73.33	100.0
W_11	398.0	292.0	195.0	209.4	85.02	100.0
W_17	392.0	289.0	194.0	208.3	75.75	100.0
W_15	381.0	285.0	198.0	210.6	71.89	100.0
W_14	376.0	277.0	190.0	203.0	69.41	100.0
W_9	361.0	269.0	184.0	197.1	69.94	100.0
W_23	367.0	265.0	173.0	188.8	69.98	100.0
W_18	361.0	261.0	169.0	184.0	72.46	100.0

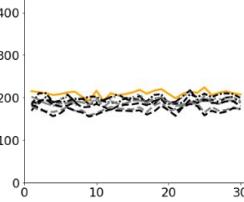
CDF OWD LATENCY (uS)



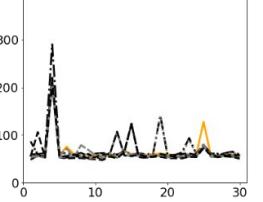
TIME SERIES P90 OWD (uS)



TIME SERIES P50 OWD (uS)



TIME SERIES STDDEV OWD (uS)

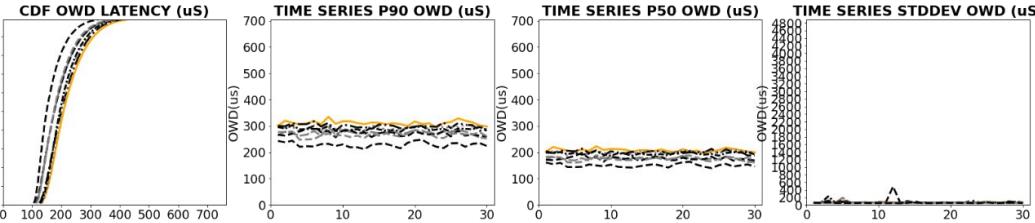


Results

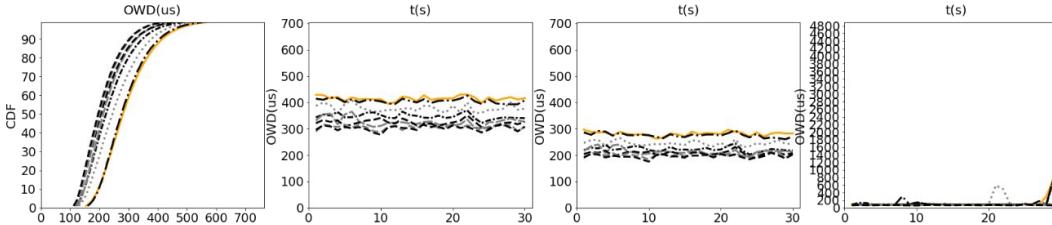
EVALUATION COMPARISON - BEST x WORST x RAND

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=30 sec

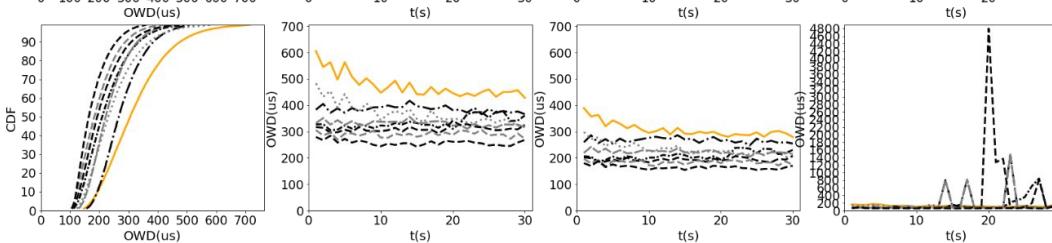
BEST-heuristic-W_3					
	99(%)	90(%)	50(%)	MEAN	STD
W_9	428.0	312.0	208.0	223.18	68.6
W_4	407.0	300.0	200.0	213.84	67.13
W_19	403.0	293.0	200.0	214.12	62.75
W_1	398.0	289.0	195.0	209.66	72.58
W_10	385.0	274.0	176.0	191.67	70.17
W_17	381.0	273.0	176.0	192.21	109.59
W_8	376.0	267.0	175.0	189.45	64.88
W_13	328.0	230.0	150.0	164.68	56.76
					100.0



WORST-p90-W_3					
	99(%)	90(%)	50(%)	MEAN	STD
W_6	568.0	415.0	281.0	299.19	172.97
W_9	554.0	406.0	276.0	294.8	173.35
W_20	526.0	378.0	245.0	262.98	164.58
W_14	491.0	348.0	223.0	240.99	164.56
W_8	456.0	327.0	214.0	228.73	77.47
W_18	460.0	323.0	206.0	223.71	159.7
W_12	424.0	307.0	201.0	214.52	68.94
W_15	421.0	303.0	196.0	209.82	70.15
					100.0



RANDxW_3					
	99(%)	90(%)	50(%)	MEAN	STD
W_20	714.0	477.0	306.0	326.82	139.82
W_22	523.0	386.0	267.0	291.76	350.35
W_21	586.0	372.0	231.0	252.98	122.74
W_16	492.0	336.0	210.0	233.77	229.18
W_19	460.0	331.0	222.0	246.79	349.25
W_1	471.0	310.0	193.0	256.17	971.11
W_14	418.0	290.0	183.0	198.06	79.52
W_18	361.0	259.0	164.0	179.22	61.21
					100.0



Results

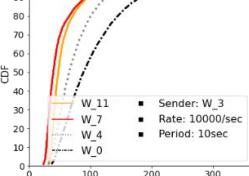
BEST HEURISTIC: STAGES 1-3

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=10 sec

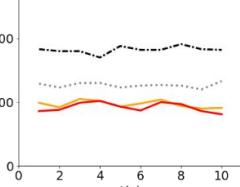
BEST-heuristic-W 3 STAGE 1

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_11	174.0	97.0	48.0	58.82	31.71	100.0
W_7	163.0	93.0	40.0	51.76	50.24	100.0
W_4	210.0	127.0	66.0	76.62	38.3	100.0
W_0	314.0	182.0	89.0	125.8	417.0	100.0

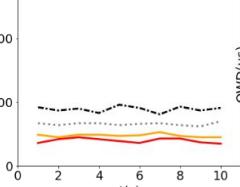
CDF OWD LATENCY (uS)



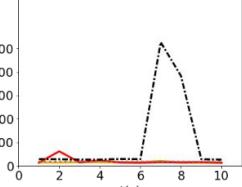
TIME SERIES P90 OWD (uS)



TIME SERIES P50 OWD (uS)



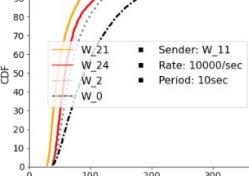
TIME SERIES STDDEV OWD (uS)



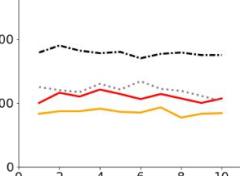
BEST-heuristic-W 3 STAGE 2

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_21	134.0	86.0	47.0	54.73	23.52	100.0
W_24	209.0	110.0	59.0	70.18	34.59	100.0
W_2	186.0	121.0	67.0	76.88	32.63	100.0
W_0	286.0	179.0	90.0	104.7	55.16	100.0

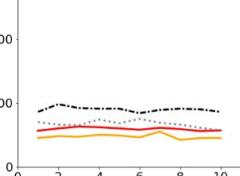
CDF OWD LATENCY (uS)



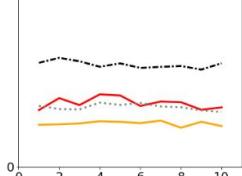
TIME SERIES P90 OWD (uS)



TIME SERIES P50 OWD (uS)



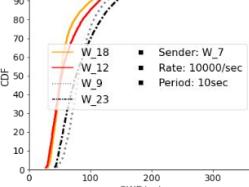
TIME SERIES STDDEV OWD (uS)



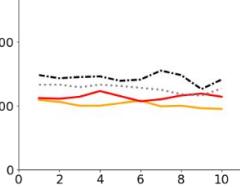
BEST-heuristic-W 3 STAGE 3

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_18	165.0	102.0	52.0	61.08	28.83	100.0
W_12	189.0	115.0	53.0	65.58	36.0	100.0
W_9	227.0	127.0	82.0	89.55	34.55	100.0
W_23	234.0	143.0	77.0	89.32	42.66	100.0

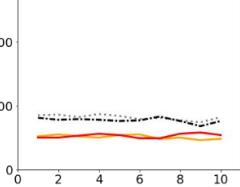
CDF OWD LATENCY (uS)



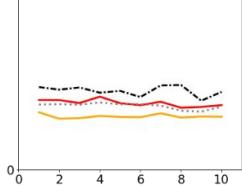
TIME SERIES P90 OWD (uS)



TIME SERIES P50 OWD (uS)



TIME SERIES STDDEV OWD (uS)



Results

BEST HEURISTIC: STAGES 4-7

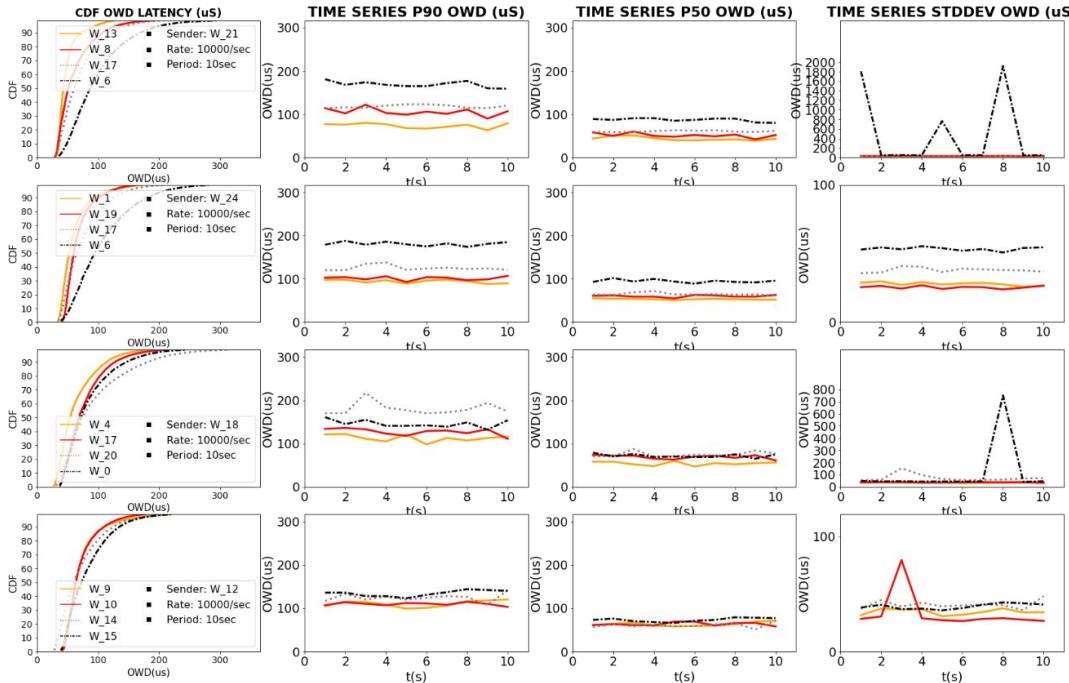
N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=10 sec

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_13	129.0	74.0	43.0	50.21	20.42	100.0
W_8	192.0	107.0	51.0	63.26	34.1	100.0
W_17	207.0	118.0	61.0	71.37	37.47	100.0
W_6	297.0	169.0	87.0	154.9	874.1	100.0

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_1	172.0	94.0	53.0	61.67	27.98	100.0
W_19	161.0	102.0	60.0	68.36	29.58	100.0
W_17	206.0	126.0	65.0	76.48	38.17	100.0
W_6	281.0	181.0	95.0	108.4	53.58	100.0

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_4	193.0	114.0	54.0	66.18	35.66	100.0
W_17	211.0	128.0	69.0	79.47	37.46	100.0
W_20	315.0	181.0	74.0	96.92	80.06	100.0
W_0	245.0	146.0	72.0	92.87	242.7	100.0

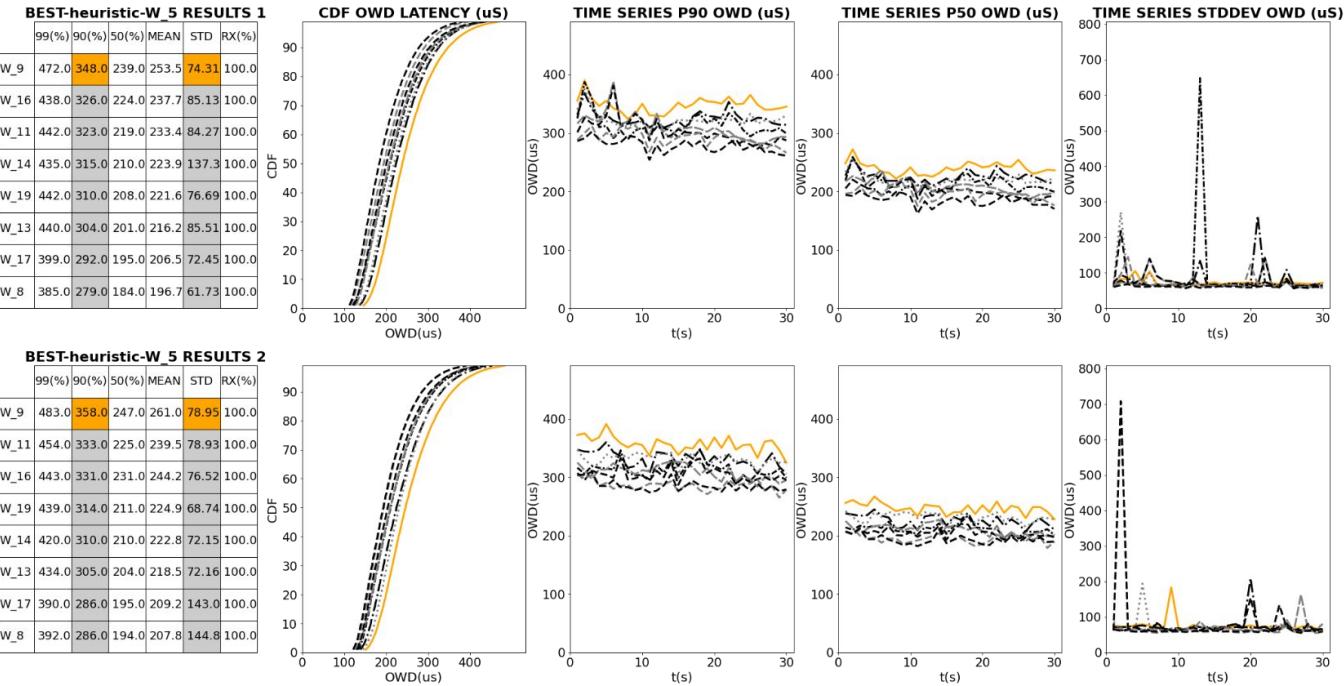
	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_9	219.0	111.0	64.0	73.63	34.85	100.0
W_10	177.0	110.0	64.0	73.11	36.84	100.0
W_14	222.0	125.0	60.0	72.53	41.23	100.0
W_15	219.0	135.0	73.0	83.8	39.51	100.0



Results

EVALUATION COMPARISON - BEST ITERATIONS

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=30 sec



- Multiple runs perform similarly

Results

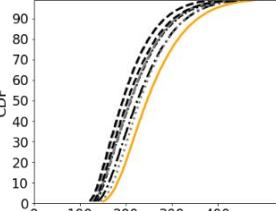
EVALUATION COMPARISON - BEST x REBUILDS

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=30 sec

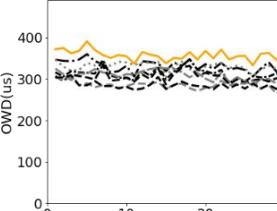
BEST-heuristic-W_5

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_9	483.0	358.0	247.0	261.0	78.95	100.0
W_11	454.0	333.0	225.0	239.5	78.93	100.0
W_16	443.0	331.0	231.0	244.2	76.52	100.0
W_19	439.0	314.0	211.0	224.9	68.74	100.0
W_14	420.0	310.0	210.0	222.8	72.15	100.0
W_13	434.0	305.0	204.0	218.5	72.16	100.0
W_17	390.0	286.0	195.0	209.2	143.0	100.0
W_8	392.0	286.0	194.0	207.8	144.8	100.0

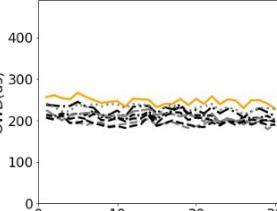
CDF OWD LATENCY (uS)



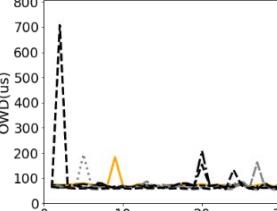
TIME SERIES P90 OWD (uS)



TIME SERIES P50 OWD (uS)



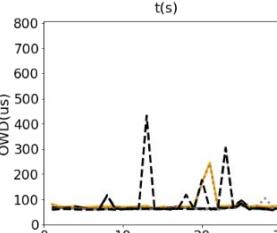
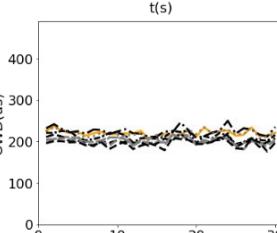
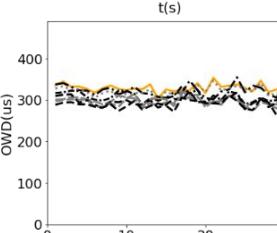
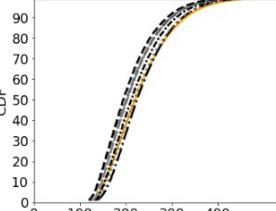
TIME SERIES STDDEV OWD (uS)



BEST-REBUILD-1-heuristic-W_5

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_12	454.0	330.0	221.0	234.4	88.58	100.0
W_16	435.0	324.0	225.0	237.0	68.87	100.0
W_11	442.0	323.0	219.0	232.8	86.16	100.0
W_19	424.0	310.0	212.0	224.5	65.09	100.0
W_13	415.0	299.0	203.0	215.8	63.53	100.0
W_14	407.0	299.0	201.0	214.2	108.5	100.0
W_17	403.0	297.0	201.0	213.3	83.73	100.0
W_8	397.0	290.0	195.0	207.9	82.78	100.0

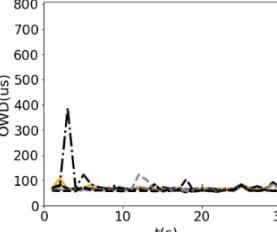
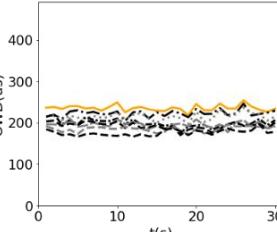
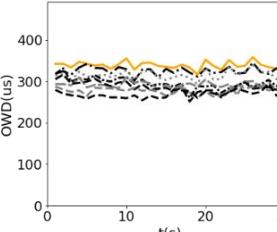
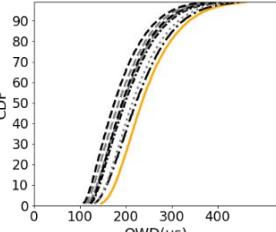
CDF



BEST-REBUILD-2-heuristic-W_5

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_22	462.0	339.0	235.0	248.8	71.55	100.0
W_16	447.0	326.0	223.0	236.1	101.3	100.0
W_11	433.0	316.0	215.0	228.0	67.82	100.0
W_19	410.0	297.0	199.0	211.9	66.22	100.0
W_14	397.0	291.0	195.0	206.7	64.04	100.0
W_13	401.0	287.0	193.0	205.7	63.6	100.0
W_17	381.0	279.0	185.0	197.2	65.66	100.0
W_8	373.0	269.0	176.0	188.6	62.43	100.0

CDF



- REBUILD on leaf layers provide inconsistent performance gains

Results

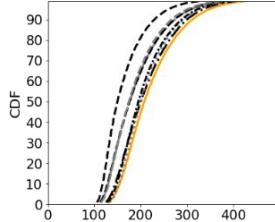
EVALUATION COMPARISON - BEST x REBUILDS

N=15, D=3, F=2, K=4, POOL=24, Rate=10000 p/sec, Duration=30 sec

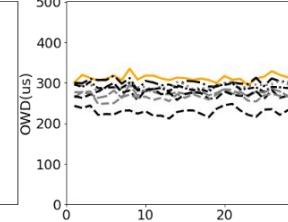
BEST-heuristic-W_3

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_9	428.0	312.0	208.0	223.1	68.6	100.0
W_4	407.0	300.0	200.0	213.8	67.13	100.0
W_19	403.0	293.0	200.0	214.1	62.75	100.0
W_1	398.0	289.0	195.0	209.6	72.58	100.0
W_10	385.0	274.0	176.0	191.6	70.17	100.0
W_17	381.0	273.0	176.0	192.2	109.5	100.0
W_8	376.0	267.0	175.0	189.4	64.88	100.0
W_13	328.0	230.0	150.0	164.6	56.76	100.0

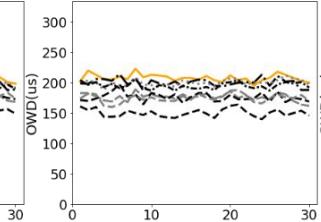
CDF OWD LATENCY (uS)



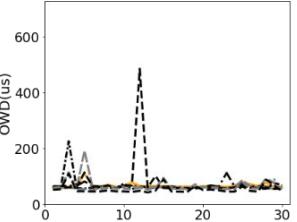
TIME SERIES P90 OWD (uS)



TIME SERIES P50 OWD (uS)



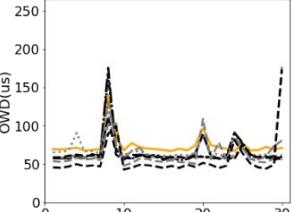
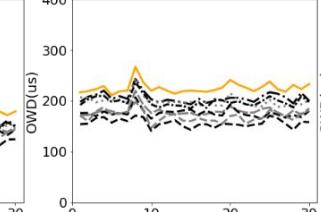
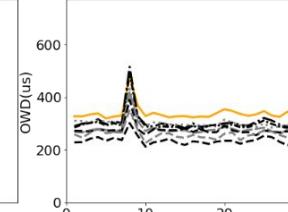
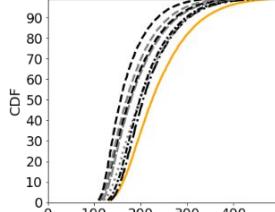
TIME SERIES STDDEV OWD (uS)



BEST-REBUILD-1-heuristic-W_3

	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_6	477.0	339.0	225.0	240.6	76.58	100.0
W_19	444.0	304.0	205.0	221.5	70.47	100.0
W_4	424.0	302.0	198.0	213.1	69.83	100.0
W_1	432.0	298.0	198.0	214.5	69.33	100.0
W_10	409.0	279.0	179.0	195.9	68.02	100.0
W_17	393.0	275.0	176.0	191.1	64.59	100.0
W_8	377.0	262.0	170.0	184.9	71.96	100.0
W_13	339.0	237.0	157.0	170.8	62.47	100.0

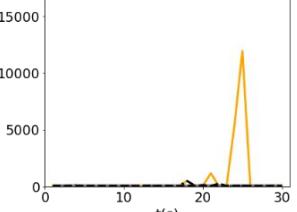
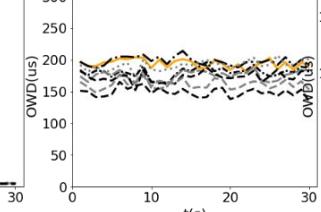
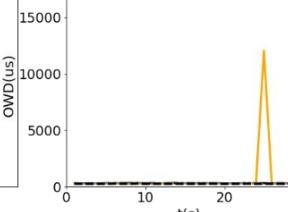
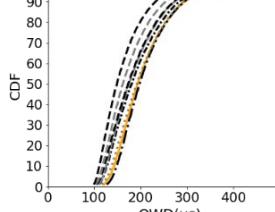
CDF



BEST-REBUILD-2-heuristic-W_3

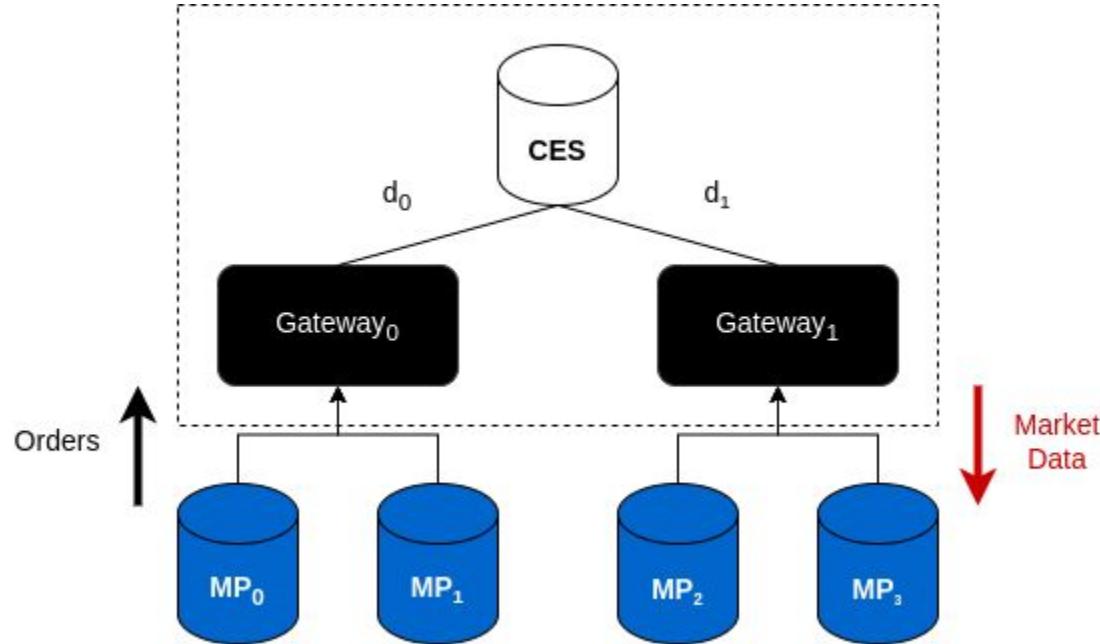
	99(%)	90(%)	50(%)	MEAN	STD	RX(%)
W_1	465.0	295.0	194.0	368.1	2507	100.0
W_19	406.0	294.0	198.0	213.1	110.1	100.0
W_4	397.0	291.0	192.0	205.2	66.06	100.0
W_2	394.0	283.0	181.0	195.5	69.08	100.0
W_10	385.0	273.0	175.0	189.6	65.98	100.0
W_17	379.0	272.0	173.0	187.5	70.49	100.0
W_8	362.0	254.0	162.0	176.9	75.72	100.0
W_13	331.0	232.0	149.0	163.1	71.2	100.0

CDF OWD (uS)



- REBUILD on leaf layers provide inconsistent performance gains

Figures 1



Figures 1